

Beyond Sequential Plans: A Three-Facet Evaluation of Agentic Planning under Plan-Preserving Perturbation

Anonymous Author(s)

Affiliation

Address

email

Abstract

A correct final answer can hide a broken plan, and a wrong final answer can hide a correct one. Evaluating LLM agents only by their final answers therefore tells us little about *where* they fail. We propose evaluating three facets of agentic behavior separately—plan structural fidelity, tool invocation accuracy, and end-task answer correctness—and build a benchmark on top of GAIA that supports all three on the same task. We add an explicit tool layer and a dependency-aware planning reference whose reordered alternatives are validated by a behavior-preserving replay filter. Across seven open-weight models spanning $50\times$ in parameter count, every single model recovers more dependencies under our augmented reference, with EdgeF1 lifting by $+45.4\%$ relative on the subset where reordering matters most—roughly three times the lift we see when this signal is averaged across all tasks. Because the three facets rank models differently—the largest open-weight model (675B) scores worse on final answers than a 31B model in the same pool—the three facets together reveal capability gaps that a single answer-rate would otherwise conflate. Code and data are available at <https://github.com/anonymoullmplanning/llm-planning>.

17 1 Introduction

Large language models are increasingly deployed as agents that decompose requests, call external tools, and synthesize final answers from intermediate observations [4, 8, 12]. Chain-of-Thought [3, 15], Tree-of-Thought [7, 17], and ReAct [11, 16] all share a common premise: before an agent executes a task, it must first form a plan that organizes subgoals, dependencies, and tool-facing actions. The plan is therefore an evaluable object, distinct from the final answer it produces.

But a benchmark that scores only the final answer cannot tell whether a failure came from a broken plan, a malformed tool call, or downstream answer synthesis—and a correct answer can mask a brittle plan that happened to succeed. We argue that a useful benchmark for agentic systems should evaluate at least three facets separately: (A) *Plan Structural Fidelity*, the dependency structure the model attributes to the task before execution; (B) *Tool Invocation Accuracy*, whether selected tools and grounded arguments cover the evidence-gathering actions the task requires; and (C) *End-Task Answer Correctness*, whether the submitted answer is accepted under standard evaluation.

Existing structural metrics already provide useful primitives: NODEF1 and EDGEF1 [2, 13] measure whether a predicted planning graph recovers the labeled nodes and dependencies of a gold reference. But the gold reference itself is the bottleneck: when it is a single serialized chain, incidental ordering looks like a required dependency, and plans that respect the underlying dependency structure but reorder independent branches are penalized. We address this by augmenting GAIA [6], which provides realistic multimodal tasks with verifiable final answers but primarily chain-style supervision. We add an explicit tool layer and convert each chain into a dependency DAG (directed acyclic graph) whose reordered alternatives are validated by a behavior-preserving replay filter, so that plans

38 matching any dependency-preserving reordering can receive credit. The synthesized dependencies
39 are validated by a three-annotator spot-check (98.4% edge correctness).

40 Across seven open-weight models spanning $50\times$ in parameter count, every single model recovers
41 more dependencies under our augmented reference, with EdgeF1 lifting by +45.4% relative on
42 the subset where reordering matters most—roughly three times the lift seen when the signal is
43 averaged across the full benchmark. The three facets rank models differently: the largest open-weight
44 model in our pool (675B) scores worse on final answers than a 31B model from the same pool, and
45 once tool execution is controlled, planning quality predicts answer correctness only weakly and
46 model-dependently. The takeaway is that planning, tool use, and answer correctness are related but
47 not interchangeable capabilities, and a benchmark that scores them separately reveals capability
48 differences that a single answer-rate would otherwise conflate.

49 This paper makes three contributions:

- 50 • **A three-facet evaluation framework** that jointly scores plan structural fidelity, tool invocation
51 accuracy, and end-task answer correctness on the same task instance, keeping downstream synthesis
52 separable from planning and tool-use errors.
- 53 • **A GAIA-derived benchmark** that adds an explicit tool layer and dependency-aware planning
54 reference to Native GAIA, with reordered alternatives validated by a behavior-preserving replay
55 filter, turning an answer-centric benchmark into a unified planning-execution-answer resource.
- 56 • **A systematic open-weight evaluation** showing that model rankings differ across the three facets,
57 that augmentation lifts dependency recovery for every model evaluated (most strongly on multi-
58 order-rich tasks), and that planning–answer correlation conditional on clean tool execution is mildly
59 positive but heterogeneous.

60 2 Related Work

61 **Reasoning to tool-augmented planning.** Recent work has transformed LLMs into tool-augmented
62 agents that solve requests through explicit intermediate reasoning and external actions [15–17]. Once
63 reasoning is externalized into a multi-step plan that drives tool invocation, the plan itself becomes an
64 evaluable object rather than a hidden by-product behind final task success.

65 **Benchmarks for tool use and task automation.** Benchmarks for tool use have evolved from API-
66 centric function-calling tests [9, 10] toward structured evaluations of task automation. TaskBench
67 introduces Tool Graphs to model dependency structure between tools [13]; GTA emphasizes human-
68 written queries with real deployed tools and multimodal inputs [14]; UltraTool expands evaluation to
69 the broader lifecycle of planning, creation, and usage [5]. Answer-centric benchmarks such as GAIA
70 measure end-to-end completion under realistic multimodal settings but do not isolate whether success
71 or failure originated in planning, execution, or final synthesis [6].

72 **Planning-aware evaluation and our setting.** A more recent line of work evaluates planning
73 and execution as distinct components of agentic behavior [1, 2]. Our setting differs in two ways.
74 First, we score the model’s planning graph separately from the realized tool trace while retaining
75 GAIA’s verifiable answer target, so the same record supports all three facets. Second, plan structural
76 evaluation uses a dependency-preserving augmented reference set rather than a single canonical chain,
77 so that plans reordering independent branches are credited rather than penalized. Table 1 summarizes
78 the dimensions that most directly motivate this setup.

79 3 LLM Planning Formalization

80 We formalize each agent run as three aligned objects per task instance: a planning graph G_{plan}
81 capturing intended subtask structure, an interaction trace τ recording realized tool use, and a final
82 answer y reflecting downstream synthesis. These three objects are the substrate for the three evaluation
83 facets; metrics over each object are defined in Section 4.5.

84 3.1 Planning Graph

85 Given a task with user query Q , the planner decomposes the task into a finite set of natural-language
 86 subgoals $\mathcal{P}_Q = \{p_1, p_2, \dots, p_n\}$, where each p_i is a planning intent (an evidence-acquisition,
 87 reasoning, or synthesis subgoal). The planner predicts dependency relations among these subgoals:
 88 $d_{ij} = (p_i, p_j)$ when p_i is required before p_j is well-defined or executable. The resulting planning
 89 graph is

$$G_{\text{plan}} = (\mathcal{P}_Q, \mathcal{D}_Q), \quad \mathcal{D}_Q = \{d_{ij}\} \subseteq \mathcal{P}_Q \times \mathcal{P}_Q.$$

90 Nodes are subgoal labels and edges are immediate dependency claims; the graph represents planning
 91 structure attributed to the model before tool-specific refinement, and is not the realized tool trace. A
 92 planning graph specifies a family of admissible realizations (the topological orders consistent with
 93 $\mathcal{D}_Q$) rather than a single fixed sequence. The gold side of the benchmark provides two reference views
 94 (Section 4.2): a sequential chain reference and an order-augmented dependency-DAG reference, both
 95 expressed over gold planning intent $\mathcal{P}_Q^*$. Planning metrics first align predicted and gold planning-
 96 intent labels, then score recovered nodes and dependency edges (Section 4.5).

97 3.2 Interaction Trace

98 Each task is associated with a prompt-visible tool subset $\mathcal{T}_Q \subseteq \mathcal{T}$, drawn from a shared runtime tool
 99 universe $\mathcal{T}$; tools outside $\mathcal{T}_Q$ are not visible to the planner. The mechanism by which $\mathcal{T}_Q$ is derived is
 100 benchmark-specific (Section 4.2). Execution is represented by an interaction trace

$$\tau = ((c_1, o_1), (c_2, o_2), \dots, (c_T, o_T)), \quad c_t = (\text{name}_t, \text{args}_t), \quad \text{name}_t \in \mathcal{T}_Q,$$

101 where name_t denotes the selected tool, args_t its grounded argument values, and o_t the resulting
 102 observation (returned result, retrieved content, or execution-time error). Not every node in G_{plan}
 103 appears in τ : some nodes correspond to internal reasoning or intermediate computation. The trace
 104 records only externally realized interactions.

105 3.3 Final Answer

106 The model submits a final answer y after conditioning on the query and accumulated execution
 107 context: $y = f_{\text{answer}}(Q, \tau)$. Final-answer evaluation requires a verifiable ground-truth answer; among
 108 existing agentic-planning benchmarks, GAIA is one of the few that supports it, and our framework
 109 therefore reports Facet C only on the GAIA-derived instantiation.

110 A complete notation reference is provided in Appendix A.1.

111 4 Benchmark Overview

112 4.1 Pipeline

113 Figure 1 summarizes our evaluation framework, which decomposes agentic-planning evaluation into
 114 three separable facets. The framework can be instantiated on any benchmark whose annotations
 115 support the corresponding facet; we additionally apply it to TaskBench and UltraTool (Appendix G)
 116 on the facets each supports. In the GAIA-derived instantiation released here, original tasks are
 117 enriched with an explicit tool layer and a dependency-aware planning reference, then processed

Benchmark	Plan	Graph	Tool-Use	Eval. Executed	Trace Realistic	Tasks Component	Eval. Dependency/Order	Refs. Flexible-Order	Scoring
TaskBench [13]	✓	✓	✓	×	×	✓	✓	×	×
Advancing Agentic Systems [2]	✓	✓	✓	✓	×	✓	✓	×	×
GAIA [6]	×	×	×	×	✓	×	×	×	×
GTA [14]	×	✓	✓	✓	✓	×	×	×	×
UltraTool [5]	✓	✓	✓	×	✓	✓	×	×	×
OrchestrationBench [1]	✓	✓	✓	✓	✓	✓	✓	×	×
Our framework	✓	✓	✓	✓	✓	✓	✓	✓	✓

Table 1: Comparison of representative planning and tool-use benchmarks. Dependency/order references indicate benchmark support for structured dependency or ordering supervision, while flexible-order scoring requires explicitly crediting dependency-preserving reorderings. Our framework is the only setup that supports flexible-order scoring against a non-canonical reference set.




Figure 1: Overall framework. *Left:* Native GAIA tasks (four domains) plus three added annotation layers form the shared record. *Middle:* A three-stage inference pipeline produces $\hat{G}_{\text{plan}}$, $\hat{\tau}$, and y (Section 3). *Right:* Three evaluation facets score plan structural fidelity (A), tool invocation accuracy (B), and end-task answer correctness (C). Five-stage planning decomposition in Appendix A.2.

through planning, tool execution, and final-answer generation, yielding the three evaluable objects scored by Facets A–C.

4.2 Augmented GAIA Construction

Our benchmark is built on GAIA, but converts GAIA from a final-answer benchmark into a resource for unified evaluation of planning, execution, and answering. We start from 165 multimodal GAIA tasks, partitioned by input modality into Text, Document, Vision, and Audio, and leave the original task statement, attachments, native chain, and answer target unchanged. Augmented GAIA adds two evaluation-facing annotation families on top of these native records: a tool-reference layer for Facet B and a dependency-aware planning layer for Facet A (Table 2).

The tool-reference layer exposes a shared 16-tool runtime universe $\mathcal{T}$ to every English GAIA task and adds 379 gold tool-slot annotations. Native GAIA does not specify a prompt-visible tool inventory or annotate intermediate tool slots, so this layer makes tool invocation directly evaluable against a stated reference rather than only reconstructable from execution traces. The same catalog is visible across domains; domain-specific demand is induced by attachments and gold references, with the detailed tool distribution reported in Appendix B.

The dependency-aware layer converts GAIA’s chain-style substeps into a gold dependency DAG for each task and derives admissible async orderings from those DAGs, with the synthesis procedure detailed in Section 4.3. The native chain remains a valid reference, while replay-filtered async orderings make dependency-preserving alternatives explicit without changing the query or answer target. After Gemma 4 behavior-preserving replay filtering, 1,357 of 1,771 candidate async orderings are retained; 58 of 165 tasks (35.2%) admit at least one retained async ordering beyond the native chain, and 54 tasks (32.7%) are multi-order-rich with at least two retained async orderings. Appendices C and E give the domain-level accounting and replay outcome breakdown. Together, these layers convert GAIA into a unified planning/execution/answer benchmark whose three facets are evaluable on the same task instance.

4.3 Dependency-Aware Data Synthesis

The original GAIA annotations provide executable solution chains but do not distinguish necessary precedence constraints from incidental serialization; we therefore construct a dependency-aware async layer from each chain. For each task Q , we treat the ordered GAIA substeps as the gold planning-intent inventory $\mathcal{P}_Q^*$ and synthesize a dependency set $\mathcal{D}_Q^*$, yielding

$$G_{\text{plan}}^*(Q) = (\mathcal{P}_Q^*, \mathcal{D}_Q^*).$$

The synthesis begins from the full ordered executable step sequence and does not apply pre-analysis pruning in the canonical construction path. This choice is deliberate: observation-like actions such as

Layer	Added Structure / Statistics	Role
Native task layer	165 GAIA tasks across Text, Document, Vision, and Audio; original query, attachments, native chain, and final-answer target retained	Raw multimodal task
Tool layer	Shared 16-tool runtime universe $\mathcal{T}$ exposed to every English GAIA task; 379 gold tool slots across 165 tasks, dominated by web search (51.7%) and web browsing (15.8%)	Enables Facet B
Planning-reference layer	Gold dependency DAG for every task; 1,284 nodes and 1,086 edges overall, averaging 7.78 nodes and 6.58 dependency edges per task	Enables Facet A
Order-augmented layer	Async order augmentation that preserves the native chain while adding dependency-preserving ordering alternatives when available; 1,771 candidate async orderings are filtered to 1,357 orderings whose <i>Gemma 4</i> replay outcome reproduces the native chain's; 58 tasks (35.2%) receive at least one retained async ordering, and 54 tasks (32.7%) are multi-order-rich	

Table 2: Augmented GAIA construction layered on top of native GAIA. Each row adds one annotation layer; rows are listed in dependency order from raw task to admissible reorderings. Quantities reported are for the final validated release after behavior-preserving replay filtering (Section 4.4).

reading, scrolling, or recording intermediate evidence may appear locally redundant, yet still establish information or state preconditions required by later actions. Pruning such steps before dependency analysis would therefore risk deleting valid precedence constraints.

Direct dependency annotation. For each task, an LLM annotator is given the task query together with the original ordered executable step sequence and is asked to identify, for each step, which earlier steps are directly required. The objective is to recover direct dependencies between steps rather than restate the full serialized chain. The resulting annotations are discarded if they do not define a valid dependency graph over the original step sequence, for example because a step is missing, a dependency points to the same or a later step, or the resulting graph contains a cycle.

We then remove redundant direct edges that are already implied by intermediate dependencies, yielding a more compact graph without changing the underlying precedence structure. The resulting graph is treated as the gold dependency DAG $G_{\text{plan}}^*(Q)$.

Admissible ordering generation. The synthesized DAG admits a family of legal topological realizations rather than a single canonical chain. We therefore derive admissible orderings that respect the dependencies in $\mathcal{D}_Q^*$, together with Kahn-level frontiers that expose coarse parallel structure. Tasks whose dependency graph admits only one linearization yield a single ordering, whereas tasks with parallel structure may yield multiple valid realizations. Because every retained dependency is forward with respect to the original executable chain, the native chain remains a valid topological realization of the synthesized graph.

This async layer is a controlled semantic extension of GAIA’s original supervision: it preserves the task, answer target, and executable step inventory, while making dependency-preserving reorderings explicit beyond a single serialized reference. Detailed synthesis and validation procedures are provided in Appendix J, while prompt contracts are summarized in Appendix K.

4.4 Dataset Validation

We validate the constructed annotations along three layers: schema and structural sanity of the LLM-synthesized dependency DAGs, behavior-preserving replay consistency of the admissible orderings under the shared executable tool layer, and a human spot-check of dependency-edge correctness.

Schema and structural validation. LLM-annotator outputs are checked for schema conformance, forward-only edge orientation under the original step order, single occurrence per step, and acyclicity; payloads failing any check are rejected. Surviving DAGs are passed through transitive reduction (Section 4.3) to produce the gold dependency DAG used by the planning-reference layer. This stage removes structural artifacts but does not validate execution feasibility.

Behavior-preserving replay filter. We additionally filter admissible orderings by whether reordering preserves the native chain’s executable behavior under a shared runtime. Each candidate ordering is replayed end-to-end with *Gemma 4* and retained when its reordered replay reproduces the native chain’s outcome (both gold-equivalent correct, or both same-wrong); otherwise it is filtered. This

operationalizes augmentation through outcome reproduction rather than outcome correctness, so the filter does not select for orderings that improve solvability. Of 1,771 candidate async orderings spanning 165 tasks, 1,357 are retained (76.6%), covering 58 tasks. Per-domain and per-difficulty breakdowns and the full outcome typology are reported in Appendix E.

Human spot-check. Three independent annotators reviewed a stratified sample of 30 multi-order-rich tasks (351 rows per annotator) to assess binary dependency quality: whether each LLM-asserted edge corresponds to a necessary parent–child dependency, and whether sampled non-edges reveal omitted necessary dependencies. On the 248 asserted edges, 244 (98.4%) were judged correct by majority vote, indicating that the synthesized graphs recover necessary dependencies rather than merely reproducing incidental chain order. On the 103 sampled non-edges, 96 (93.2%) were judged to correctly reflect the absence of a dependency, while 7 (6.8%) were judged as omitted dependencies; agreement on these non-edge judgments is strong (Fleiss $\kappa = 0.755$). Free-form annotations identify a small number of additional omitted dependencies. Per-annotator agreement, the prevalence-deflated Fleiss κ for asserted edges, and the full omitted-dependency breakdown are reported in Appendix F.

4.5 Evaluation Metrics

We define metrics over the three evaluation objects (Section 3): *PlanningScore* on G_{plan} , *ToolUsageScore* on τ , and Exact Match (EM), Token F1, and a manually-adjusted *Accepted Answer* rate on y . We additionally define two execution-side rates—*Aligned Tool Success* (ATS) and *Tool Invocation Success Rate* (TISR)—used as filtering criteria for the *tool-clean subset* (Section 5), not as model-comparison metrics.

Planning metrics. Planning metrics compare the predicted planning graph $\hat{G}_{\text{plan}}$ against the gold reference set, which contains the native chain reference and the admissible orderings synthesized in Section 4.3. For each predicted graph, we compute the planning score against every reference in the set and report the maximum, so that a prediction matching any valid ordering is credited.

NodeF1 measures node-side recovery after one-to-one semantic alignment. Let M be the maximum-cardinality alignment between predicted intents $\hat{p} \in \hat{\mathcal{P}}$ and gold intents $p \in \mathcal{P}^*$ whose label similarity $s(p, \hat{p})$ (sentence-embedding cosine) exceeds threshold ξ ($\xi = 0.45$). Then $\text{Precision}_{\text{node}} = |M|/|\hat{\mathcal{P}}|$, $\text{Recall}_{\text{node}} = |M|/|\mathcal{P}^*|$, and $\text{NodeF1} = F_1(\text{Precision}_{\text{node}}, \text{Recall}_{\text{node}})$.

EdgeF1 measures direct dependency recovery after the same node alignment. Each predicted dependency $\hat{d}_{ij} = (\hat{p}_i, \hat{p}_j) \in \hat{\mathcal{D}}$ is mapped into gold-intent space when both endpoints are aligned by M . Let $\tilde{\mathcal{D}}$ be the resulting aligned predicted dependency set. Then $\text{Precision}_{\text{edge}} = |\tilde{\mathcal{D}} \cap \mathcal{D}^*|/|\tilde{\mathcal{D}}|$, $\text{Recall}_{\text{edge}} = |\tilde{\mathcal{D}} \cap \mathcal{D}^*|/|\mathcal{D}^*|$, and $\text{EdgeF1} = F_1(\text{Precision}_{\text{edge}}, \text{Recall}_{\text{edge}})$.

The headline planning score is $\text{PlanningScore} = \alpha \text{NodeF1} + (1 - \alpha) \text{EdgeF1}$, with $\alpha = 0.5$ as default. A soft node-label diagnostic *NodeLabelSim* used internally by the node matcher is defined in Appendix I.

Tool usage and execution metrics. Tool-usage metrics operate on realized tool calls in $\hat{\tau}$, namely $c_t = (\text{name}_t, \text{args}_t)$. The headline *ToolUsageScore* averages three components: tool-family selection (ToolF1), argument-slot prediction (ParamF1), and argument-value grounding (ParamValueF1), each defined as F_1 of the corresponding precision and recall. The default headline score is $\text{ToolUsageScore} = (\text{ToolF1} + \text{ParamF1} + \text{ParamValueF1})/3$; alternative weightings are discussed in Appendix I.

We define two execution-side rates used to construct the tool-clean subset. Let $s_t \in \{0, 1\}$ indicate whether invocation $\hat{c}_t$ executed successfully, g_ℓ be the number of gold tool slots requiring tool family ℓ , and $\hat{s}_\ell$ be the number of successful predicted calls using tool family ℓ . Then

$$\text{TISR} = \frac{1}{\hat{T}} \sum_{t=1}^{\hat{T}} s_t, \quad \text{ATS} = \frac{\sum_{\ell} \min(g_\ell, \hat{s}_\ell)}{\sum_{\ell} g_\ell}.$$

TISR captures per-task tool-execution stability (whether predicted calls actually run), while ATS captures whether successfully-executed calls cover the gold tool families. The tool-clean subset (Appendix H) is the set of (task, model) pairs with no parse error, $\text{ATS} \geq 0.5$, and $\text{TISR} \geq 0.7$.

Final answer metrics. Final-answer metrics operate on the submitted answer y . We report Exact Match (EM) and Token F1 against the gold answer. EM is strict and penalizes superficial surface variations such as plural-singular mismatches, trailing punctuation, or whitespace differences that do not affect answer correctness. To recover such cases, we additionally report an *Accepted Answer* rate: every model output that fails strict EM is manually inspected and accepted if it matches the gold answer modulo surface normalization. The Accepted Answer rate is the headline final-answer metric used throughout the experiments.

5 Experiments

5.1 Experimental Settings

We evaluate the following open-weight model pool on all 165 tasks in the GAIA-derived benchmark: *Mistral Large 3 675B*, *Llama-3.1 405B*, *Llama-3.3 70B*, *Gemma 4 31B*, *Mistral Small 3.2 24B*, *Llama-4 Maverick 17B*, and *Gemma 3 12B*. The pool spans more than $50\times$ in parameter count and includes models from four families, allowing us to examine whether the augmentation effect generalizes across model architectures rather than reflecting an artifact of any single training pipeline. A closed-model extension covering eight API-dependent OpenAI- and Gemini-family models is reported in Appendix L; the augmentation lift remains positive for all eight models.

We focus the main-text experiments on the GAIA-derived benchmark, which is the only dataset in our suite that supports all three facets jointly. Cross-benchmark validation on TaskBench and UltraTool—which support Facet A (planning) and Facet B (tool usage) but lack the verifiable final answers required for Facet C—is reported in Appendix G under the same framework.

Inference setup. For every (model, task) pair, planning is scored on the model’s predicted planning graph, tool usage and execution are scored on the realized trace, and final-answer correctness is scored on the submitted response. All runs share the same executable tool environment (Section 4.2) and the same metric definitions (Section 4.5). Additional implementation details are reported in Appendix A, bootstrap procedures in Appendix A.3, and prompt contracts in Appendix K.

For vision tasks, models access visual evidence through the shared image-recognition tool rather than receiving raw images in the planner prompt, preserving a uniform tool-invocation interface across models with heterogeneous multimodal capability; the vision-tool backend is described in Appendix K.7.

Augmented reference construction. Unless otherwise stated, *Native GAIA* uses the original serialized chain reference and *Augmented GAIA* adds the behavior-preserving replay-filtered async layer (Section 4.4); tasks with no retained async ordering fall back to the native chain alone.

Tool-clean filter. The tool-clean filter (Section 4.5) is applied only to analyses that ask whether planning quality predicts final-answer correctness after the tool layer is controlled (Table 5). The full three-facet headline (Section 5.2) and the multi-order score-lift analysis (Section 5.3) use the full GAIA task pool unless explicitly stated; ATS and TISR definitions, threshold sensitivity testing, and the tool-intensive variant are reported in Appendix H.

5.2 Overall Three-Facet Performance on Augmented GAIA

Table 3 reports headline performance across the three facets. Three observations stand out.

(i) The three facets do not collapse into a single capability ranking. *Gemma 4* (31B) attains the highest Accepted Answer rate at 0.424, more than double ($\sim 2.2\times$) the rate of *Mistral Large 3* (675B) at 0.194, despite the larger model recovering comparable planning nodes (NodeF1 0.514 vs. 0.519). Final-answer correctness is therefore not predictable from parameter count alone, nor from any single facet, illustrating the value of reporting all three facets jointly rather than collapsing them into a summary score.

(ii) Augmentation lifts EdgeF1 but leaves NodeF1 unchanged. The augmented reference set raises mean EdgeF1 from 0.095 to 0.107 (+12.6% relative), with all listed models showing a positive lift; NodeF1 is identical under both references by construction. The headline lift is partly diluted

Model	Params	(A) Plan Structural Fidelity			(B) Tool Invocation Accuracy			(C) End-Task Answer Correctness			
		NODEF1 (Native = Aug.)	EDGEF1		TOOLF1	PARAMF1	PARAMVALUEF1	EM	TOKENF1	Raw rate	Accepted*
			Native	Aug. Δ							
<i>Mistral Large 3</i>	675B	0.514	0.101	0.107 +0.006	0.412	0.458	0.206	0.182	0.223	0.224	0.194
<i>Llama-3.1</i>	405B	0.519	0.107	0.118 +0.011	0.473	0.518	0.279	0.176	0.217	0.261	0.182
<i>Llama-3.3</i>	70B	0.486	0.091	0.105 +0.014	0.527	0.572	0.232	0.182	0.232	0.219	0.206
<i>Gemma 4</i>	31B	0.519	0.101	0.116 +0.015	0.537	0.579	0.264	0.339	0.384	0.445	0.424
<i>Mistral Small 3.2</i>	24B	0.512	0.096	0.109 +0.013	0.545	0.569	0.233	0.103	0.148	0.179	0.145
<i>Llama-4 Maverick</i>	17B	0.494	0.092	0.105 +0.013	0.424	0.454	0.201	0.218	0.263	0.294	0.261
<i>Gemma 3</i>	12B	0.499	0.079	0.091 +0.012	0.549	0.567	0.241	0.091	0.136	0.158	0.115
Model mean	—	0.506	0.095	0.107 +0.012	0.495	0.531	0.237	0.184	0.229	0.254	0.218

Table 3: Three-facet headline results on Augmented GAIA over $N = 165$ tasks per model (1,155 total model-task predictions). *Plan Structural Fidelity*: NodeF1 is invariant to reordering by construction (identical under native and augmented references); EdgeF1 is reported under both references with Δ showing the augmentation lift. *Tool Invocation Accuracy*: ToolF1 / ParamF1 / ParamValueF1 measure correct tool selection, argument-slot prediction, and argument-value grounding (Section 4.5). *End-Task Answer Correctness*: EM and TokenF1 are automatic; *Raw rate* is the unprocessed correctness ratio; The *Accepted** column is the manually adjusted correctness rate after surface-form review and is the headline final-answer metric used throughout the experiments.

Subset (n)	Metric	Native	Augmented	Δ	Δ (%)	95% CI of Δ	Models with $\Delta > 0$
GAIA Level 2* ($n = 29$)	NODEF1	0.539	0.539	+0.000	+0.0%	— (by construction)	0/7
	EDGEF1	0.088	0.124	+0.036	+40.6%	[+0.0271, +0.0439]	7/7
	PLANNINGSCORE	0.313	0.331	+0.018	+5.7%	[+0.0135, +0.0219]	7/7
All* ($n = 54$)	NODEF1	0.515	0.515	+0.000	+0.0%	— (by construction)	0/7
	EDGEF1	0.074	0.107	+0.033	+45.4%	[+0.0263, +0.0398]	7/7
	PLANNINGSCORE	0.294	0.311	+0.017	+5.7%	[+0.0132, +0.0199]	7/7

* Multi-order-rich denotes tasks with at least two retained async orderings beyond the native chain.

Table 4: Score lift on the replay-filtered multi-order-rich subset, where the augmented reference set differs from the native chain. Each row reports the mean score across the 7-model pool, with bootstrap 95% CI on Δ from 1,000 resamples. EdgeF1 lifts by $\sim 45.4\%$ relative under augmentation on the full multi-order-rich slice, with all 7 models showing a positive shift on both subsets, while NodeF1 is unchanged by construction. The relative lift on the multi-order-rich subset is roughly three times the headline lift in Table 3 (+12.6% on the full benchmark), confirming that augmentation’s effective magnitude is partly diluted on the full benchmark by tasks with a single admissible ordering.

281 because most tasks have a single admissible ordering and therefore an identical reference under both
282 conditions; on the subset of tasks where the augmented reference differs from the native chain, the
283 effect is substantially larger (Section 5.3).

284 (iii) **Argument-value grounding remains the weakest tool-side signal.** ParamValueF1 (mean 0.237)
285 is consistently lower than ToolF1 (mean 0.495) and ParamF1 (mean 0.531) for every model: models
286 select the right tool family more reliably than they ground its argument values, locating the strongest
287 tool-side gap at argument grounding rather than tool selection.

288 5.3 Score Lift on the Multi-Order-Rich Subset

289 We now isolate tasks where the synthesized DAG admits at least two retained async orderings beyond
290 the native chain, yielding 54 tasks (39 *Text*, 9 *Document*, 5 *Vision*, 1 *Audio*; 29 at GAIA Level 2). On
291 this subset the augmented reference contains at least three orderings, so the augmentation effect is
292 exposed by a non-trivial alternative set.

293 Table 4 reports the score lift. EdgeF1 rises by +0.033 (+45.4% relative) on the full multi-order-rich
294 subset and +0.036 (+40.6% relative) on the Level 2 slice—roughly three times the headline lift
295 in Table 3—with all 7 models positive on both subsets. On the multi-order-rich subset, 19.0% of
296 open-weight task-model pairs are best matched by an augmented ordering rather than the native
297 chain, so the lift is concentrated in genuinely parallel tasks rather than spread uniformly across
298 pairs. NodeF1 is unchanged by construction, confirming that augmentation operates on dependency
299 structure: the same planning-intent inventory is preserved while incidental serialization is relaxed
300 into validated alternatives. A qualitative example illustrating this effect is given in Appendix D.

Model	Params	N	Accepted rate	r_{native}	r_{aug}	Δr	95% CI of Δr
<i>Mistral Large 3</i>	675B	53	0.170	−0.156	−0.141	+0.014	[−0.025, +0.074]
<i>Llama-3.1</i>	405B	49	0.184	+0.217	+0.207	−0.010	[−0.034, +0.000]
<i>Llama-3.3</i>	70B	58	0.172	−0.080	−0.037	+0.043	[−0.022, +0.145]
<i>Gemma 4</i>	31B	62	0.371	+0.258	+0.276	+0.018	[−0.007, +0.053]
<i>Mistral Small 3.2</i>	24B	60	0.150	+0.039	+0.056	+0.016	[−0.039, +0.092]
<i>Llama-4 Maverick</i>	17B	59	0.254	+0.045	+0.079	+0.034	[−0.027, +0.110]
<i>Gemma 3</i>	12B	57	0.088	+0.042	+0.027	−0.015	[−0.046, −0.000]
Model mean / total	–	398	0.201	+0.052	+0.066	+0.014	–

Table 5: Per-model task-level Pearson correlation between PlanningScore and Accepted Answer rate on the GAIA Level 2 tool-clean *Text/Document/Audio* subset. N is the per-model tool-clean task count; r_{native} and r_{aug} are correlations under native and augmented references respectively. Bootstrap 95% CI on Δr from 1,000 paired resamples. The macro mean Δr is +0.014, with 5/7 models showing a positive shift; several model-level CIs include zero. The result is read as a within-model diagnostic rather than a cross-model claim (see Section 5.4 for interpretation).

5.4 Planning–Answer Correlation on Tool-Clean Level 2 Tasks

To examine whether planning quality predicts final-answer correctness after tool execution is controlled, we compute task-level Pearson correlation within each model’s tool-clean subset (no parse error, $\text{ATS} \geq 0.5$, $\text{TISR} \geq 0.7$) on GAIA Level 2 *Text/Document/Audio* tasks (Table 5); Level 1 and Level 3 are excluded for insufficient length and answer sparsity, *Vision* for small task count.

The result is a within-model diagnostic. The direction is mildly positive on average ($\bar{\Delta r} = +0.014$, with 5/7 models showing a positive shift), but heterogeneous, and several model-level CIs include zero. The augmentation only changes the reference set on the multi-order-rich subset (54 of 165 tasks), so Δr measured on the full Level 2 tool-clean sample is necessarily diluted by tasks where native and augmented references coincide. The link between structural planning quality and final-answer correctness, once tool execution is controlled, is therefore best read per model rather than aggregated across the pool.

6 Limitations

The framework and its GAIA-derived instantiation have several boundaries. First, our augmentation reduces false negatives from surface-form deviation but does not detect false positives: a plan that violates true dependencies but surface-matches a reference can still receive high EdgeF1. Second, the behavior-preserving replay filter operationalizes plan equivalence through outcome reproduction under a single replay model (*Gemma 4*): an ordering is retained when its reordered replay reproduces the native chain’s outcome, whether correct or same-wrong. Two consequences follow. The augmentation effect we measure is restricted to orderings whose reordered execution behavior is reproducible by this one model, and orderings whose reordered replay yields a different outcome are filtered out even when the underlying plan structure is dependency-preserving. Third, statistical inference is constrained on multiple axes: 165 tasks across imbalanced domains (*Vision 10*, *Audio 3* are reported descriptively), a 7-model pool with limited cross-model bootstrap power, and a planning–answer correlation diagnostic (Table 5) computed on the full tool-clean Level 2 sample rather than the multi-order-rich subset, diluting Δr . Finally, the joint three-facet evaluation requires verifiable answers and is currently demonstrated only on GAIA; cross-benchmark validation on TaskBench and UltraTool (Appendix G) covers Facet A and Facet B alone.

7 Conclusion & Future Work

This paper argues that realistic agent evaluation should not collapse planning, execution, and final answering into a single terminal score. Starting from GAIA, we add an explicit tool layer and a dependency-aware planning reference whose reordered alternatives are validated by a behavior-preserving replay filter, so that the same task can be evaluated at the level of the planning graph, the realized interaction trace, and the submitted answer.

The empirical results show why this separation matters. Across seven open-weight models spanning $50\times$ in parameter count, every model recovers more dependencies under the augmented reference

(+45.4% relative EdgeF1 on multi-order-rich tasks). The three facets rank models differently—675B scores worse on final answers than 31B—so planning, tool grounding, and answering are related but not interchangeable, and scoring them separately reveals gaps a single answer-rate would conflate.

The closed-model extension (Appendix L) confirms the augmentation lift across all 15 models, but the planning–answer correlation remains heterogeneous and would benefit from broader model-family coverage. Separately, EDGEF1 rewards local dependency recovery and can understate structurally valid but differently shaped plans; richer graph-level metrics are a natural next step.

References

- [1] A. Ahn, S. Lee, H. Wang, C. Park, D. Kim, J. Roh, K. Yang, W. Jang, H. Woosung, M. S. Kim, and J. Kang. Orchestrationbench: Llm-driven agentic planning and tool use in multi-domain scenarios. In *International Conference on Learning Representations*, 2026.
- [2] A. G. Gabriel, A. A. Ahmad, and S. K. Jeyakumar. Advancing agentic systems: Dynamic task decomposition, tool integration and evaluation using novel metrics and dataset. *arXiv preprint arXiv:2410.22457*, 2024.
- [3] S. Gao, J. Dwivedi-Yu, P. Yu, X. E. Tan, R. Pasunuru, O. Golovneva, K. Sinha, A. Celikyilmaz, A. Bosselut, and T. Wang. Efficient tool use with chain-of-abstraction reasoning. In *Proceedings of the 31st International Conference on Computational Linguistics*, 2025.
- [4] S. Hong, M. Zhuge, J. Chen, X. Zheng, Y. Cheng, J. Wang, C. Zhang, Z. Wang, S. K. S. Yau, Z. Lin, L. Zhou, C. Ran, L. Xiao, C. Wu, and J. Schmidhuber. MetaGPT: Meta programming for a multi-agent collaborative framework. In *International Conference on Learning Representations*, 2024.
- [5] S. Huang, W. Zhong, J. Lu, Q. Zhu, J. Gao, W. Liu, Y. Hou, X. Zeng, Y. Wang, L. Shang, X. Jiang, R. Xu, and Q. Liu. Planning, creation, usage: Benchmarking llms for comprehensive tool utilization in real-world complex scenarios. *arXiv preprint arXiv:2401.17167*, 2024.
- [6] G. Mialon, R. Dessi, M. Lomeli, C. Eichenberg, L. Bandarkar, D. Hupkes, T. Lago, L. F. R. Ribeiro, A. Auffray, S. Katsumata, et al. Gaia: A benchmark for general ai assistants. *arXiv preprint arXiv:2311.12983*, 2023.
- [7] Z. Ni, Y. Li, N. Yang, D. Shen, P. Lyu, and D. Dong. Tree-of-code: A self-growing tree framework for end-to-end code generation and execution in complex tasks. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025.
- [8] J. S. Park, J. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023.
- [9] S. G. Patil, T. Zhang, X. Wang, and J. E. Gonzalez. Gorilla: Large language model connected with massive apis. *arXiv preprint arXiv:2305.15334*, 2023.
- [10] Y. Qin, S. Liang, Y. Ye, K. Zhu, L. Yan, Y. Lu, Y. Lin, X. Cong, X. Tang, B. Qian, S. Zhao, L. Hong, R. Tian, R. Xie, J. Zhou, M. Gerstein, D. Li, Z. Liu, and M. Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis. *arXiv preprint arXiv:2307.16789*, 2023.
- [11] J. Shen, H. Bai, L. Zhang, Y. Zhou, A. Setlur, S. Tong, D. Caples, N. Jiang, T. Zhang, A. Talwalkar, and A. Kumar. Thinking vs. doing: Improving agent reasoning by scaling test-time interaction. In *Advances in Neural Information Processing Systems*, 2026.
- [12] Y. Shen, K. Song, X. Tan, D. Li, W. Lu, and Y. Zhuang. HuggingGPT: Solving AI tasks with chatGPT and its friends in hugging face. In *Advances in Neural Information Processing Systems*, 2023.
- [13] Y. Shen, K. Song, X. Tan, W. Zhang, K. Ren, S. Yuan, W. Lu, D. Li, and Y. Zhuang. Taskbench: Benchmarking large language models for task automation. *arXiv preprint arXiv:2311.18760*, 2024.
- [14] J. Wang, Z. Ma, Y. Li, S. Zhang, C. Chen, K. Chen, and X. Le. Gta: A benchmark for general tool agents. *arXiv preprint arXiv:2407.08713*, 2024.

- 387 [15] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou.
388 Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural*
389 *Information Processing Systems*, 2022.
- 390 [16] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. React: Synergizing
391 reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.
- 392 [17] S. Yao, D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan. Tree of thoughts:
393 Deliberate problem solving with large language models. In *Advances in Neural Information*
394 *Processing Systems*, 2024.

Technical Appendices

The appendices provide audit material for the main claims. Appendix A reports implementation details (notation, pipeline, bootstrap, environment, architecture). Appendices B–F document the dataset: tool-demand distribution (B), augmented GAIA data profile (C), a qualitative multi-order example (D), behavior-preserving replay filter (E), and human dependency spot check (F). Appendix G validates the framework on TaskBench and UltraTool. Appendix H details the tool-clean filtering variants. Appendix I defines metric-weighting conventions and diagnostic components. Appendix J gives the full dependency-synthesis procedure. Appendix K reproduces prompt contracts and templates. Appendix L extends the evaluation to eight API-dependent closed models. Unless otherwise stated, each table reports the evaluation slice or construction pool named in its caption.

A Experiments Implementation Details

A.1 Notation Convention

Throughout the paper, the superscript $\star$ denotes a gold reference and the hat $\hat{\cdot}$ denotes a model prediction; for example, $G_{\text{plan}}^\star$ and $\hat{G}_{\text{plan}}$ refer to the gold and predicted planning graphs respectively, and analogously for $\mathcal{P}_Q^\star$, $\hat{\mathcal{P}}_Q$, $\mathcal{D}_Q^\star$, $\hat{\mathcal{D}}_Q$, and τ , $\hat{\tau}$. When a quantity refers to both interchangeably, we drop the decoration. Tool spaces are denoted at two levels: $\mathcal{T}$ is the shared runtime universe across the benchmark, and $\mathcal{T}_Q \subseteq \mathcal{T}$ is the prompt-visible subset exposed to the planner for task Q .

A.2 Pipeline Implementation Detail

The three-stage inference pipeline summarized in Figure 1 is implemented as a five-stage procedure for reproducibility. The planning stage is decomposed into three substages—abstract planning (the model decomposes the query into natural-language subgoals), tool instantiation (each subgoal is grounded to one or more candidate tool invocations from $\mathcal{T}_Q$), and plan refinement (the candidate set is consolidated into a coherent executable plan). The execution stage corresponds to iterative tool calls and observation collection along the realized trace $\hat{\tau}$, and the answer stage produces the final response y . The five-stage decomposition is internal to the planning sub-procedure and does not change the evaluable objects ($\hat{G}_{\text{plan}}$, $\hat{\tau}$, y); the three-stage framing in the main text suffices for evaluation purposes.

A.3 Bootstrap Details

The main score-lift table uses a cross-model bootstrap over per-model score deltas. Model indices are resampled with replacement for 1,000 iterations using seed 42, and the reported interval is the percentile 95% confidence interval of the mean per-model delta. The task pool is the Gemma 4 behavior-preserving replay-filtered multi-order-rich subset: 54 tasks total, including 29 GAIA Level 2 tasks, with no tool-clean filter. NODEF1 has zero delta because native and augmented references use the same planning-intent inventory.

The main planning–answer correlation table uses per-model paired task bootstrap. For each model, retained task rows are resampled with replacement using the same sampled row indices for native PLANNINGSCORE, augmented PLANNINGSCORE, and Accepted Answer. Degenerate bootstrap samples with constant planning score or constant Accepted Answer are skipped. The tool-clean filter is no parse error, $\text{ATS} \geq 0.5$, and $\text{TISR} \geq 0.7$.

A.4 Environment Setup

Experiments are run in a Python 3.10 environment on Ubuntu 22.04.5. The local machine has two Intel Xeon Gold 6430 processors (128 logical cores), 503 GiB RAM, and two NVIDIA RTX PRO 6000 Blackwell GPUs with 96 GB VRAM each. Local compute is used for embedding-based scoring, file processing, and tool-side execution when available.

Model generations are served through externally supplied API endpoints, including OpenAI-compatible wrappers where used. API credentials are provided at runtime and are not stored in the

benchmark data or paper artifacts. No benchmarked model is trained or fine-tuned; all reported model outputs are generated in inference-only mode.

GAIA, TaskBench, and UltraTool evaluations share the same evaluation code path and result schema, while dataset-specific records determine attachments, gold tool references, and final-answer availability. The main GAIA evaluation covers 165 tasks across seven open-weight models, yielding 1,155 model–task runs; the closed-model extension adds eight API-dependent models, yielding 1,320 additional model–task runs. Runtime varies with endpoint latency, tool execution, and max-turn usage; recomputing metrics from cached outputs is substantially faster and can be done on CPU-only workers.

A.5 Existing Assets and Terms

Table 6 summarizes the external assets used in the study and how they are credited. We do not redistribute third-party model weights. The released benchmark artifacts consist of derived annotations, evaluation code, normalized result summaries, Croissant metadata, and rebuild scripts. They do not redistribute GAIA raw questions, final answers, or attachments; users remain responsible for obtaining source benchmarks through their official access channels and following the corresponding licenses and provider terms.

Asset	Source / version used	License or terms treatment
GAIA	GAIA validation tasks and attachments from the official benchmark release [6]; dataset source: https://huggingface.co/datasets/gaia-benchmark/GAIA .	GAIA is distributed through a gated Hugging Face dataset card that asks users not to reshare the validation or test set in crawlable form; those source terms remain in force. Our release therefore provides derived dependency, tool-reference, and replay-filtered ordering annotations plus rebuild scripts, but does not redistribute GAIA raw questions, final answers, or attachments.
TaskBench	Public TaskBench benchmark used for cross-benchmark Facets A–B [13]; dataset source: https://huggingface.co/datasets/microsoft/Taskbench .	MIT license. We report evaluation results and do not change the source benchmark license.
UltraTool	Public UltraTool benchmark used for cross-benchmark Facets A–B [5]; source repository: https://github.com/JoeYing1019/UltraTool .	Apache License 2.0. We report evaluation results and do not change the source benchmark license.
Evaluated model families	Llama, Gemma, Mistral, OpenAI-family, and Gemini-family models accessed through externally supplied API-compatible endpoints.	Model generations are used under the applicable Meta Llama, Google Gemma, Mistral AI, OpenAI API, Google Gemini API, or endpoint-provider terms. No model weights are redistributed and no model is trained or fine-tuned.
Scoring dependencies	Python tooling and the sentence-transformers/all-MiniLM-L6-v2 embedding model used for node-label matching.	Used as runtime dependencies under their respective package/model licenses; third-party dependency code is not republished as part of the benchmark artifact.

Table 6: Existing assets used by the benchmark and evaluation.

A.6 Architecture Details

Planning is evaluated on the abstract planning-intent graph emitted by the model. Native GAIA uses the original chain reference, while Augmented GAIA scores the same planning-intent inventory against the dependency-aware reference set after behavior-preserving replay filtering. Node label matching uses sentence-transformers/all-MiniLM-L6-v2 with node alignment threshold $\xi = 0.45$. The data synthesis and dependency annotator is *GPT-4o*. Tool usage is evaluated from the realized tool trace, and final-answer correctness is evaluated from the submitted answer field used for human review.

B Tool-Demand Distribution Across GAIA Tasks

The main experimental settings in Section 5 define the model pool, decoding setup, and tool-clean filter. Here we document the evidence-gathering tool demand of the 165-task GAIA set, excluding the framework-side final-answer submission action. Figure 2 visualizes the within-split gold tool-slot distribution. For *Document*, attachment-access steps are displayed by their typed document-reader realization—spreadsheet, generic file, PDF, presentation, or archive reader—while the scoring layer also accepts the original code-execution realization as an equivalent way to access the same document evidence.




Figure 2: Within-split gold tool-slot distribution across the 165-task GAIA set. Bars are normalized within each split and show relative tool composition rather than raw counts. In *Document*, attachment-access calls are displayed as typed document-reader calls: 29.2% spreadsheet reader, 10.4% generic file reader, 6.3% PDF reader, 4.2% archive reader, and 2.1% presentation reader.

473 The split-level profile is intentionally skewed. *Text* is search-heavy, *Document* concentrates on typed
 474 attachment access plus reasoning, *Vision* combines image recognition with calculation, browsing,
 475 and code execution, and *Audio* pairs transcription with downstream reasoning.

Domain	Main tool families used by the gold references	Role in the benchmark
<i>Text</i>	web search, web browsing, calculator, reasoning, Python execution	Search and inspect sources, then compute intermediate quantities.
<i>Document</i>	spreadsheet reader, generic file reader, PDF reader, archive extractor, presentation reader, reasoning	Parse local attachments and combine document evidence with reasoning.
<i>Vision</i>	image recognition, calculator, web browsing, Python execution, web search	Convert visual evidence into textual or structured observations before computing or inspecting supporting sources.
<i>Audio</i>	audio transcription, reasoning	Transcribe audio evidence and solve text constraints over the transcript.

Table 7: Domain-level tool demand under the shared 16-tool executable catalog. All English GAIA tasks see the same 16 prompt-visible tools; domain-specific demand is reflected in task attachments and gold tool references rather than in a different visible tool list.

476 C Augmented GAIA Data Profile

477 C.1 Dataset Views and Accounting

478 Augmented GAIA keeps the original 165 GAIA tasks, queries, attachments, and final-answer targets
 479 unchanged. The release adds three evaluation-facing layers: a tool-reference layer for grounded
 480 tool scoring, *GPT-4o* dependency annotations for planning-reference construction, and a Gemma
 481 4 behavior-preserving replay-filtered async-ordering view used as the Augmented GAIA reference
 482 ordering layer. The unfiltered dependency view contains the task-level DAGs together with all
 483 candidate dependency-preserving orderings derived from them. The filtered ordering view keeps
 484 the same task-level dependency DAGs but retains only behavior-preserving async orderings under
 485 the Gemma 4 replay policy. The native chain remains available for every task, so tasks with no
 486 retained async ordering fall back to the original chain reference rather than being removed from the
 487 benchmark.

Quantity	Value	Interpretation
Tasks	165	Original GAIA task instances retained without changing query, attachment, or answer target.
Planning nodes	1,284	Step inventory in the <i>GPT-4o</i> dependency DAG annotations.
Dependency edges	1,086	Reduced dependency edges used by the planning-reference layer.
Candidate async orderings	1,771	Candidate dependency-preserving non-native orderings generated from the unfiltered dependency view.
Retained async orderings	1,357	Gemma 4 behavior-preserving non-native orderings retained in the filtered ordering view.
Native-only tasks	107	Tasks with no retained async ordering; the native chain remains the sole reference.
Tasks with ≥ 1 retained async ordering	58 (35.2%)	Tasks with at least one retained reference beyond the native chain.
Tasks with exactly 1 retained async ordering	4 (2.4%)	Native chain plus one retained async reference.
Tasks with ≥ 2 retained async orderings	54 (32.7%)	Multi-order-rich subset used for flexible-order stress tests.
Mean native steps per task	7.78	Average size of the original step inventory.
Mean dependency edges per task	6.58	Average reduced dependency-edge count.

Table 8: DAG and paper-accounting summary computed directly from the final Augmented GAIA release files.

488 C.2 Replay-Filtered Coverage




Figure 3: Replay-filtered Augmented GAIA coverage. The task bar shows the share of tasks with at least one retained async ordering beyond the native chain; the ordering bar shows the Gemma 4 behavior-preserving retention rate over candidate async orderings.




Figure 4: Augmented GAIA construction profile by input domain and GAIA level. Stacked bars decompose each slice’s tasks into native-only, exactly-one-retained, and multi-order-rich cases; the right columns report task counts, retained candidate-ordering rows, retention rate, and mean native steps/dependency edges per task, with the final row giving the full-release total.

489 This profile separates three notions that are easy to conflate. First, the dependency DAG is task-level
490 structure: all 165 tasks receive a *GPT-4o* dependency DAG over the original planning-intent steps.
491 Second, candidate async orderings are construction-time realizations sampled from those DAGs; they
492 are not new tasks and do not change the answer target. Third, Gemma 4 replay filtering selects the
493 ordering rows that preserve native-chain execution behavior under the shared tool layer. The result is
494 a conservative augmentation: 107 tasks remain native-chain only, 4 tasks add exactly one retained
495 async reference, and 54 tasks form the multi-order-rich subset used for the strongest flexible-order
496 diagnostics.

C.3 Reference-Validated Critical Path vs. Native Chain Length

To check whether the final Augmented GAIA reference set exposes meaningful parallelism, we summarize the replay-filtered reference orderings rather than unfiltered dependency annotations. For each task q , let L_q be the native chain length. We induce a conservative partial order from the final Augmented GAIA references by keeping an ordering constraint only when it appears in every retained reference ordering for that task; let C_q be the longest path length in this reference-validated partial order. We define the parallelism ratio as

$$\rho_q = 1 - \frac{C_q}{L_q}.$$

A value of $\rho_q = 0$ means the final references preserve the native chain order, while larger values indicate replay-validated reordering structure. Normalized Kendall distance is row-level and measures how far each retained reference ordering moves from the native chain. Table 9 first summarizes the three replay-filtered coverage buckets with task-weighted critical-path statistics, and Table 10 reports the row-level reordering diagnostic over increasingly strict retained-ordering thresholds. The ≥ 2 threshold is the multi-order-rich subset used in the main flexible-order diagnostics.

Task bucket	Tasks	Native mean	Native median	Critical mean	Critical median	Parallelism mean
All tasks	165	7.78	7.0	6.45	6.0	11.6%
Native-only	107	6.35	6.0	6.35	6.0	0.0%
Exactly-one	4	5.50	5.5	4.50	4.5	19.0%
Multi-order-rich	54	10.80	9.0	6.81	6.5	33.9%

Table 9: Task-weighted reference-validated critical-path profile by replay-filtered coverage bucket. Critical path is computed from the conservative partial order induced by final Augmented GAIA references, so native-only tasks preserve the native chain by definition.

Threshold	Tasks	Retained rows	Parallelism mean / med.	Kendall dist. mean / med. / max
≥ 1	58	1,357	40.0% / 40.0%	0.230 / 0.200 / 0.667
≥ 2 (multi-order-rich)	54	1,353	40.1% / 40.0%	0.230 / 0.200 / 0.667
≥ 5	40	1,314	40.4% / 40.0%	0.232 / 0.205 / 0.667
≥ 10	32	1,266	41.0% / 40.0%	0.236 / 0.209 / 0.667
≥ 20	26	1,163	42.2% / 42.9%	0.246 / 0.222 / 0.667
≥ 50	13	650	42.6% / 40.0%	0.259 / 0.222 / 0.667

Table 10: Reference-validated critical-path and reordering diagnostic over Gemma 4-retained non-native async-ordering thresholds. Parallelism is computed from the final Augmented GAIA reference orderings; Kendall distance is computed between each retained reference ordering and the native chain.

510 D Qualitative multi-order example: Native GAIA chain vs. Augmented 511 GAIA DAG




Figure 5: Qualitative multi-order example from a Document task. The native chain serializes the two accommodation-type computations, while the augmented DAG exposes them as independent branches before the final comparison.

512 This example illustrates the kind of planning behavior that the augmented reference is intended to
 513 credit. The task asks the model to compare the average ratings of two accommodation types in an
 514 attached PDF. The native chain is a valid execution: open the PDF, compute the hotel total and mean,
 515 then compute the rental-house total and mean, and finally compare the two averages. However, the
 516 ordering between the hotel branch and the rental-house branch is incidental to the solution. Once the
 517 PDF has been read, the two branch-level summaries can be computed independently before the final
 518 comparison.

519 The augmented dependency DAG preserves the same step inventory and answer target but removes
 520 the requirement that the hotel branch must precede the rental-house branch. Both branches remain
 521 dependent on opening the PDF, and the final compare node depends on both branch means. Thus,
 522 the augmented reference does not change the task or invent an alternative solution; it exposes the
 523 partial order that is already implicit in the native solution. This is the distinction that EDGEF1 is
 524 meant to measure: whether the predicted planning graph recovers necessary dependencies rather than
 525 the arbitrary serialization of one recorded execution.

526 The model plan on the right captures the same planning intent. The raw emitted abstract plan
 527 extracts the accommodations and ratings, groups them by type, computes the average rating for
 528 each accommodation type, and compares the averages. In the figure, the average-computation node
 529 is expanded into hotel and rental-house mean branches to make the recoverable dependency fork
 530 explicit. Under the native sequential chain, this intent can be penalized because it does not reproduce
 531 the hotel-then-rental order exactly. Under Augmented GAIA, the same plan receives credit because it
 532 respects the dependency DAG: evidence extraction precedes grouping, the two means are independent
 533 given the grouping, and the answer depends on their comparison. In this instance, the model is
 534 *Llama-3.1 405B*; its native PLANNINGSORE is 0.525, while its augmented PLANNINGSORE is
 535 0.622, with paper EDGEF1 increasing from 0.250 to 0.444 (+0.194 absolute, +77.8% relative using
 536 the unrounded scores) and the final stage matching the gold answer *Hotels*.

E Behavior-Preserving Replay Filter for Augmented GAIA Reference Construction

The async ordering layer of Augmented GAIA is filtered by a behavior-preserving replay procedure under *Gemma 4* as the replay model. The filter operationalizes plan equivalence through outcome reproduction (Section 4.4): a candidate ordering is retained when its reordered replay reproduces the native chain’s final-answer outcome on the same task under the shared executable tool layer, and is otherwise filtered out. This frames augmentation as a conservative extension of the native chain rather than as a filter that selects only rows solved correctly by a strong model.

Outcome reproduction is operationalized as gold-evaluator equivalence on the answer string. Two replays reproduce the same outcome in either of two cases: (i) both replays produce gold-equivalent correct answers, possibly with surface-form differences (numeric formatting such as 1 vs. 1.0, accepted aliases such as St. Petersburg vs. Saint Petersburg, or tolerance-equivalent numbers such as 563.9 vs. 564.0); or (ii) both replays produce the same wrong answer. Orderings whose reordered replay yields a different outcome than the native chain—changing a correct answer to a wrong one, recovering a correct answer where the native chain failed, or producing a different wrong answer—are filtered. Retaining the same-wrong category is the conservative choice for plan equivalence: when the native and reordered replays converge to the same incorrect outcome, the reordering has not altered the executable behavior of the native plan, even if that behavior happens to be incorrect on the gold target. The filter does not select for orderings that improve solvability.

Table 11 summarizes the filter result. Of 1,771 candidate async orderings, 1,357 (76.6%) are retained and 414 (23.4%) are filtered. Among the 414 filtered orderings, 88 (21.3% of filtered) have no usable native-chain replay output, so the row is not treated as an augmentation-specific failure: the native chain itself does not produce a definite outcome to reproduce. The remaining 326 (78.7% of filtered) correspond to ordering-specific non-reproduction, where the native chain produces an outcome but the reordered replay produces a different one.

Item	Count	Rate
Total candidate ordering rows	1,771	100.0% (1,771/1,771)
Retained under behavior-preserving filter	1,357	76.6% (1,357/1,771)
Filtered under behavior-preserving filter	414	23.4% (414/1,771)
Native chain replay unavailable	88	21.3% of filtered (88/414)
Ordering-specific non-reproducible outcome	326	78.7% of filtered (326/414)

Table 11: *Gemma 4* behavior-preserving replay filter summary for Augmented GAIA reference construction.

Table 12 reports the full outcome typology for all 1,771 candidate orderings. The first three rows are retained under the behavior-preserving policy (1,219 both-correct text-equivalent, 37 both-correct surface-different, and 101 both-wrong same-answer); the remaining five rows are filtered. The two largest filtered categories are *both wrong, different wrong answers* (151, 8.5% of all candidates) and *native correct, reordered wrong* (69, 3.9%). The latter is the augmentation-specific failure mode where reordering an originally correct execution produces a wrong outcome under *Gemma 4*; the former represents executions that fail in different ways under reordering, indicating that the reordered trajectory diverges from the native one.

Replay outcome	Count	Rate
Native and reordered answers both correct and text-equivalent	1,219	68.8% (1,219/1,771)
Native and reordered answers both correct but surface-different	37	2.1% (37/1,771)
Both answers wrong, same wrong answer	101	5.7% (101/1,771)
Both answers wrong, different wrong answers	151	8.5% (151/1,771)
Native correct, reordered wrong	69	3.9% (69/1,771)
Native wrong, reordered correct	68	3.8% (68/1,771)
Native chain replay produced no usable answer	88	5.0% (88/1,771)
Reordered replay produced no usable answer	38	2.1% (38/1,771)

Table 12: *Gemma 4* replay outcome typology over all 1,771 candidate async orderings. Rows 1–3 are retained under the behavior-preserving filter; rows 4–8 are filtered.

Two design points are worth noting for reproducibility. First, the filter operationalizes plan equivalence through outcome reproduction under a single replay model, so the retention pool reflects *Gemma 4*’s execution profile in particular; an alternative multi-model agreement filter would yield a smaller pool with stronger cross-model behavior alignment, at the cost of further reducing the retained orderings beyond the present 1,357. Second, retaining *both wrong, same wrong answer* rows is a deliberate conservatism: these orderings preserve native-chain executable behavior under reordering, even though the underlying behavior is itself incorrect on the gold target, and excluding them would convert the filter into a correctness-conditional selector rather than a behavior-preservation filter. Section 6 discusses the implications of these design choices for the scope of the augmentation effect we report.

F Human Dependency Spot Check

To assess the semantic quality of the synthesized dependency graphs, we conduct an independent human spot check over 30 tasks. The sampled tasks are selected from instances with multiple admissible execution orders, where dependency quality is especially consequential: a spurious edge can unnecessarily constrain the execution graph, while a missing edge can admit invalid reorderings.

Three independent annotators review each sampled task. For each task, annotators are given the original query, the chain-style step descriptions, the synthesized dependency edges, and a sampled set of step pairs for which no edge was asserted. Annotators are instructed to judge dependencies using only the provided query and step sequence, without external search or additional task execution. This design isolates whether the dependency relation is supported by the annotated executable plan itself.

The annotation form contains two row types. For *asserted edges*, annotators judge whether the proposed parent–child dependency actually exists. An asserted edge is judged correct if the parent step provides information, state, context, or an intermediate result required for the child step to be validly executed. These rows estimate the false-positive rate of the synthesized dependency graph. For *sampled nonedges*, annotators judge whether the absence of an edge is correct or whether a necessary dependency was missed. These rows estimate false negatives among sampled unasserted step pairs. A separate free-form field allows annotators to record additional missing dependencies outside the sampled nonedge rows.

The completed forms contain 351 judgments per annotator: 248 asserted edges and 103 sampled nonedges. Pooling both row types as binary dependency-existence judgments, annotators are unanimous on 330/351 rows (94.0%), with Fleiss $\kappa = 0.905$. Pairwise Cohen κ values are 0.959, 0.885, and 0.872. This high pooled agreement indicates that the binary distinction between dependency and non-dependency is stable under independent human review.

For asserted edges, majority vote judges 244/248 edges correct (98.4%) and 4/248 incorrect (1.6%). Annotators are unanimous on 232/248 asserted edges (93.5%). Fleiss κ for asserted-edge correctness is 0.363, with pairwise Cohen κ values of 0.605, 0.293, and 0.274. This lower chance-corrected agreement should be interpreted together with the strong prevalence imbalance in this subset: because nearly all asserted edges are judged correct, the expected chance agreement is high, compressing κ despite high raw agreement.

For sampled nonedges, majority vote judges 96/103 pairs as correctly having no dependency (93.2%) and 7/103 as missed dependencies (6.8%). Annotators are unanimous on 98/103 sampled nonedges (95.1%). Fleiss κ is 0.755, with pairwise Cohen κ values of 0.928, 0.693, and 0.641. The free-form fields identify 10 unique additional missing dependencies; 8 receive two-annotator support, while 2 receive single-annotator support.

Overall, the spot check indicates that the synthesized dependency graphs have high precision on asserted edges and limited false negatives among sampled nonedges. The remaining errors are concentrated in a small number of omitted dependencies rather than widespread spurious dependencies. These results support the use of the synthesized graphs as binary dependency references for order-aware planning evaluation.

Model	Params	(A) Plan Structural Fidelity			(B) Tool Invocation Accuracy			(C) End-Task Answer Correctness			
		NODEF1	EDGEF1	PLANNINGSCORE	TOOLF1	PARAMF1	PARAMVALUEF1	EM	TOKENF1	Raw rate	Accepted rate
TaskBench											
Mistral Large 3	675B	0.668	0.235	0.451	0.659	0.457	0.266	–	–	–	–
Llama-3.1	405B	0.771	0.367	0.569	0.767	0.600	0.393	–	–	–	–
Llama-3.3	70B	0.705	0.337	0.521	0.725	0.537	0.341	–	–	–	–
Gemma 4	31B	0.839	0.439	0.639	0.822	0.655	0.453	–	–	–	–
Mistral Small 3.2	24B	0.792	0.398	0.595	0.751	0.596	0.387	–	–	–	–
Llama-4 Maverick	17B	0.812	0.436	0.624	0.774	0.610	0.406	–	–	–	–
Gemma 3	12B	0.723	0.318	0.521	0.708	0.525	0.331	–	–	–	–
Model mean	–	0.759	0.361	0.560	0.744	0.569	0.368	–	–	–	–
UltraTool											
Mistral Large 3	675B	0.492	0.079	0.286	0.866	0.630	0.395	–	–	–	–
Llama-3.1	405B	0.517	0.105	0.311	0.932	0.682	0.380	–	–	–	–
Llama-3.3	70B	0.492	0.103	0.298	0.740	0.518	0.312	–	–	–	–
Gemma 4	31B	0.558	0.116	0.337	0.943	0.795	0.549	–	–	–	–
Mistral Small 3.2	24B	0.524	0.107	0.315	0.916	0.670	0.412	–	–	–	–
Llama-4 Maverick	17B	0.518	0.120	0.319	0.955	0.709	0.445	–	–	–	–
Gemma 3	12B	0.516	0.107	0.312	0.850	0.622	0.366	–	–	–	–
Model mean	–	0.517	0.105	0.311	0.886	0.661	0.408	–	–	–	–

Table 13: Cross-benchmark results on TaskBench and UltraTool. Facet C is blank because neither dataset provides verifiable final-answer targets.

G Cross-Benchmark Scope

We use TaskBench and UltraTool to test whether the end-to-end evaluation machinery transfers to public benchmarks beyond our GAIA-derived setting. *Augmented GAIA* is the benchmark instantiation in this study that fits the full three-facet framework: native GAIA supplies verifiable final-answer targets, while our annotations add the explicit tool layer and dependency-aware planning references needed for planning and tool-use evaluation.

TaskBench and UltraTool are used as subsampled cross-benchmark checks with the same seven-model pool used in the main GAIA tables. For TaskBench, we construct a balanced 1,000-example subset with 500 DAG-style examples and 500 chain-style examples. For UltraTool, we use a 1,000-example English subset. To fit these datasets into our staged evaluation framework, we normalize their original records into a shared evaluation format, including task records, planning references, tool-reference fields, and model-facing input templates. These format adjustments are used only to make the benchmarks compatible with our evaluation pipeline and do not change the underlying task intent.

Both datasets support the planning and tool-use facets, but they lack the human-auditable final-answer target used for the Accepted Answer rate. Facet C is therefore intentionally left blank for these benchmarks. Their role is to test whether the planning/tool evaluation machinery transfers beyond *Augmented GAIA*, rather than to extend final-answer evaluation.

H Tool-clean Variants

H.1 Tool-Clean Filtering Rules

A prediction is *tool-clean* when it has no parse error, $ATS \geq 0.5$, and $TISR \geq 0.7$. ATS is a gold-slot coverage measure: it asks how many required gold tool slots are matched by successful non-submit predicted tool calls. TISR is an execution-stability measure: it asks what fraction of the model’s non-submit tool calls completed without runtime failure. The two quantities are complementary: a model can execute many calls successfully while still missing the gold tool slots, or it can invoke a relevant tool but fail because the argument or runtime environment is invalid.

The main planning–answer correlation diagnostic uses the GAIA Level 2 *Text/Document/Audio* slice. Level 1 tasks are often too short for planning variation to be visible, Level 3 tasks are answer-sparse, and *Vision* has too few tasks for stable per-model correlation estimates. Table 14 summarizes the main tool-clean slice and the corresponding tool-intensive variant.

Slice	Extra filter	N	Accepted rate	Mean r_{native}	Mean r_{aug}	Mean Δr
Level 2 <i>Text/Document/Audio</i> tool-clean	None	398	0.201	0.050	0.065	+0.014
Tool-intensive variant	Gold tool slots ≥ 2	335	0.212	0.070	0.086	+0.016

Table 14: Main correlation slice and tool-intensive variant. Tool-intensive means tool-clean plus at least two gold tool slots. Mean Δr is averaged over per-model deltas, so it can differ slightly from subtracting the rounded mean-correlation columns.

H.2 Threshold Sensitivity

Table 15 varies the ATS and TISR thresholds while keeping the no-parse-error requirement and the same Level 2 *Text/Document/Audio* base slice fixed. These rows should be read as ATS/TISR-controlled variants of the tool-clean diagnostic; only the $\text{ATS} \geq 0.5$, $\text{TISR} \geq 0.7$ row is the main tool-clean setting used elsewhere. This main setting keeps enough samples for per-model task-level correlation while still removing obvious tool-execution failures. Stricter thresholds reduce sample size sharply; the direction remains mostly positive, but the number of positive models falls from 5/7 to 3/7 under the strictest settings.

ATS threshold	TISR threshold	N	Accepted rate	Mean r_{native}	Mean r_{aug}	Mean Δr	Positive models
0.5	0.5	422	0.201	0.037	0.051	+0.014	5/7
0.5	0.7	398	0.201	0.050	0.065	+0.014	5/7
0.5	0.9	342	0.211	0.049	0.068	+0.018	5/7
0.7	0.7	210	0.176	0.079	0.094	+0.015	3/7
0.7	0.9	180	0.167	0.072	0.094	+0.022	3/7
0.9	0.7	185	0.184	0.061	0.080	+0.019	3/7
0.9	0.9	157	0.172	0.054	0.080	+0.025	3/7

Table 15: Threshold sensitivity for the Level 2 *Text/Document/Audio* tool-clean slice.

H.3 Tool-Intensive Filtering Rules

The tool-intensive subset keeps the same tool-clean rule and additionally requires at least two gold tool slots. This isolates tasks where tool-mediated evidence gathering is more central, while preserving the same execution-control threshold used in the main correlation slice. Table 16 reports the per-model diagnostic for this subset.

Model	N	Accepted rate	r_{native}	r_{aug}	Δr	95% CI of Δr
<i>Mistral Large 3</i>	45	0.156	-0.111	-0.092	+0.020	[-0.029, +0.086]
<i>Llama-3.1</i>	43	0.186	0.200	0.189	-0.011	[-0.039, +0.000]
<i>Llama-3.3</i>	48	0.208	-0.044	-0.003	+0.041	[-0.033, +0.142]
<i>Gemma 4</i>	52	0.404	0.272	0.289	+0.017	[-0.012, +0.054]
<i>Mistral Small 3.2</i>	50	0.140	0.051	0.073	+0.022	[-0.047, +0.110]
<i>Llama-4 Maverick</i>	51	0.275	0.046	0.088	+0.042	[-0.025, +0.131]
<i>Gemma 3</i>	46	0.087	0.092	0.073	-0.019	[-0.066, +0.001]
Model mean / total	335	0.212	0.072	0.088	+0.016	—

Table 16: Tool-intensive correlation diagnostic on the GAIA Level 2 *Text/Document/Audio* slice.

I Metric Weighting and Diagnostic Components

I.1 NodeLabelSim Diagnostic

NodeLabelSim is the soft node-label diagnostic used internally by the node matcher. Let $s(p, \hat{p})$ be the sentence-embedding cosine similarity between a gold planning-intent label p and a predicted planning-intent label $\hat{p}$. We compute

$$\text{NodeLabelSim} = \frac{1}{|\mathcal{P}^*|} \sum_{p \in \mathcal{P}^*} \max_{\hat{p} \in \hat{\mathcal{P}}} s(p, \hat{p}).$$

This score is used internally by the node matcher to seed the alignment M used by NodeF1; it is not used as a headline model-comparison metric.

668 The main planning score is a weighted average of planning-intent node recovery and direct dependency
669 recovery:

$$\text{PlanningScore}(\alpha) = \alpha \text{NodeF1} + (1 - \alpha) \text{EdgeF1}.$$

670 The main text uses $\alpha = 0.5$, giving node recovery and dependency recovery equal weight. The
671 headline tool-grounding score is

$$\text{ToolUsageScore}(\beta) = \beta_{\text{tool}} \text{ToolF1} + \beta_{\text{param}} \text{ParamF1} + \beta_{\text{value}} \text{ParamValueF1},$$

672 where the nonnegative weights sum to one. Unless otherwise noted, the reported results use equal
673 weights, $\beta_{\text{tool}} = \beta_{\text{param}} = \beta_{\text{value}} = 1/3$. Table 17 summarizes how the individual components
674 should be read.

Component	What it emphasizes	Interpretation
NODEF1	Planning-intent recovery	Checks whether the model states the required planning steps after semantic one-to-one matching.
NODELABELSIM	Soft label similarity	Records embedding similarity between gold planning nodes and their nearest predicted labels before thresholded matching.
EDGEF1	Direct dependency recovery	Checks whether the recovered steps have the right immediate dependency relations.
PLANNINGSCORE	Graph recovery	Averages node recovery and direct dependency recovery with $\alpha = 0.5$ in the main tables.
TOOLF1	Tool-family selection	Rewards choosing the right tool families.
PARAMF1	Argument-slot selection	Rewards supplying the required argument fields; non-required execution-control fields are normalized before scoring.
PARAMVALUEF1	Argument-value grounding	Rewards grounding the argument values to the correct resource, query, URL, expression, or equivalent execution object.

Table 17: Metric-component interpretation.

675 I.2 ParamValueF1 Matching Details

676 For GAIA, the main PARAMVALUEF1 scorer uses execution-normalized stable value slots. Value
677 credit is considered only after a predicted tool and parameter slot align with a gold slot; the remaining
678 comparison asks whether the argument value points to the same executable object or information
679 need. This normalization resolves benign runtime-surface variation: attachment paths may match by
680 GAIA path or basename, typed file readers are mapped to equivalent file-access realizations, image
681 and audio paths are mapped to a shared file resource, URLs are normalized after removing unstable
682 query noise, calculator expressions may match by numeric equivalence, and query-like strings use
683 token and character overlap. Free-form reasoning prompts, custom image prompts, and raw code
684 text are not treated as stable value slots by default, although file resources referenced inside code are
685 still credited. The older strict value score is kept as an audit field, but the main TOOLUSAGESCORE
686 uses the execution-normalized value component. This is why PARAMVALUEF1 remains lower than
687 TOOLF1: selecting the correct tool family is often easier than grounding the exact resource, query,
688 expression, or file object used by that tool.

689 J Dependency-Synthesis Details for Augmented GAIA

690 GAIA’s native annotations record one successful execution chain per task. A chain identifies a
691 single valid serialization of the solution but does not distinguish semantically required precedences
692 from incidental ordering choices. This distinction matters for order-aware planning evaluation: a
693 reference should penalize violations of necessary dependencies without penalizing valid reorderings
694 or admissible partial parallelization.

695 To address this, we synthesize a gold dependency DAG $G_{\text{plan}}^*(Q) = (\mathcal{P}_Q^*, \mathcal{D}_Q^*)$ from each chain-style
696 annotated solution via the four-stage procedure below. The synthesis preserves the original task query,
697 gold final answer, and executable step inventory; it changes only the dependency structure imposed
698 over the recorded steps.

699 J.1 Starting Point

700 Dependency analysis is performed over the full recorded step sequence $\mathcal{P}_Q^*$ without pre-analysis
701 pruning. This is a deliberate methodological choice: observation-like steps such as reading, scrolling,

or recording intermediate evidence may appear locally redundant but can establish information or state preconditions required by later actions. Pruning such steps before annotation risks deleting valid precedence constraints.

All synthesized edges are restricted to forward directions with respect to the original chain, i.e. for any retained dependency $(p_i^*, p_j^*) \in \mathcal{D}_Q^*$ we require $i < j$. This reflects the fact that the source chain is already a successful execution; the synthesized graph explains which earlier steps are necessary for later ones rather than introducing dependencies that contradict the recorded solution trajectory.

J.2 Dependency Annotation

An LLM annotator (*GPT-4o*) is given the task query Q and the full ordered step sequence $\mathcal{P}_Q^*$, and is asked to identify, for each step, which earlier steps are *directly* required—rather than merely restating the recorded order. The annotation targets immediate parent–child dependencies only: if p_i^* enables p_k^* and p_k^* enables p_j^* , then p_i^* should not be listed as a direct parent of p_j^* unless there is also an independent direct dependency. This convention keeps the annotation focused on immediate explanatory dependencies while the subsequent reduction stage removes any remaining redundant edges.

A recurring challenge in web-grounded plans is cross-branch reuse: a later action may depend on a value, URL component, or identifier first obtained in a different part of the execution. Such dependencies can be visually distant in the recorded chain. The annotation process therefore treats cross-branch reuse as a first-class dependency source rather than a local surface pattern.

J.3 Well-Formedness Validation

Each candidate annotation is checked against three structural requirements before being accepted: (i) every step in $\mathcal{P}_Q^*$ appears exactly once as a child, with root steps allowed an empty parent set; (ii) every proposed parent of step p_j^* is a valid earlier step p_i^* with $i < j$; and (iii) no duplicate parent assignments exist for the same child. Annotations violating any condition are discarded rather than heuristically repaired, preventing malformed annotations from entering the benchmark.

The validated annotation is then converted into a directed graph and checked for acyclicity. Any cyclic candidate is discarded. This check is necessary because the annotator reasons locally over natural-language descriptions, and local mutual-necessity judgments can otherwise produce cycles.

J.4 Transitive Reduction

The validated graph is simplified by transitive reduction: any edge (p_i^*, p_j^*) that is already implied by an intermediate directed path of length at least two is removed. Reachability witnesses are evaluated against the full validated graph rather than a progressively reduced one, making the retained edge set independent of removal order. The result is the final gold dependency DAG $G_{\text{plan}}^*(Q) = (\mathcal{P}_Q^*, \mathcal{D}_Q^*)$, which is more compact than the validated candidate graph while preserving the same partial order.

J.5 Admissible Realizations

$G_{\text{plan}}^*(Q)$ defines a partial order over $\mathcal{P}_Q^*$. Its admissible realizations are all topological orders consistent with $\mathcal{D}_Q^*$. Because every retained edge satisfies $i < j$, the native chain $\mathcal{P}_Q^*$ is itself one valid topological realization; the synthesis therefore embeds the original chain inside a larger family of dependency-respecting orderings rather than replacing it.

For benchmark construction we derive two summaries from $G_{\text{plan}}^*(Q)$. A *level decomposition* is obtained by repeatedly extracting the zero-indegree frontier, exposing coarse parallel structure. A *reproducible realization set* enumerates a deterministic, task-local subset of admissible topological orders: when the number of linear extensions is small this coincides with the complete set; otherwise a bounded sampling procedure is used.

K Prompt Contracts and Templates

This appendix summarizes the prompt contracts that define the evaluation-facing behavior of the staged planner. Rather than reproducing every instruction in full, we describe the stable interface at each stage: the information shown to the model, the decision requested, and the structured response expected. Run-specific details such as formatting reminders, language-specific suffixes, attachment listings, and dynamic execution notes are omitted because they vary by task but do not change the underlying contracts.

These stages instantiate the three-stage pipeline in Figure 1: Stages 1–3 implement planning, Stage 4 produces the realized interaction trace, and Stage 5 produces the final answer. Facet A uses the Stage-1 $\hat{G}_{\text{plan}}$ by default. When an evaluated run has an empty Stage-1 planning graph, Stage-3 is used as a planning-metric fallback only; tool traces and submitted answers are not replaced.

Stage	Prompt-visible input	Required model behavior
System JSON guard	Provider-specific system role or equivalent instruction	Return a single JSON object with no prose, markdown fence, or hidden thinking tags. This guard is prepended to API calls and is supplemented by stage-local output schemas.
Abstract planning	User request and optional record-level extra instruction; no concrete tool list	Produce task-specific natural-language subgoals that form $\hat{G}_{\text{plan}}$. Edges must express true information or state dependencies rather than the incidental order of a chain. This graph is the planning-intent object used by Facet A.
Tool sufficiency	User request, abstract plan, and benchmark-visible tools	Decide whether existing visible tools can execute the abstract plan. New tools are allowed only for a clear capability gap; the default expected output is an empty new-tool list.
Concrete execution planning	User request, attachments, abstract plan, visible executable tools, and any executable Stage-2 handoff	Map the abstract plan into a concrete execution DAG and planned tool calls. File tools must receive the exact attachment path; audio files should start with transcription; final-answer submission is excluded from Stage 3 planning nodes.
Iterative execution	Task, visible runtime tools, Stage-3 plan, prior tool calls, observations, and dynamic attachment/reflection guidance	Choose exactly one next action per turn. The model must return a tool choice, arguments, and brief reasoning; if the answer is known, it should submit the final answer. Repair prompts reuse this same one-action contract.
Final answer synthesis	User request and bounded execution observations	Produce a structured final-answer object for pre-final synthesis and audit context. The reviewed answer remains the explicit Stage-4 submitted answer.
Vision tool backend	Image file, task string, optional OCR/color candidates, and a structured extractor system instruction	Extract visual evidence for the planner rather than solve the task directly. Task variants cover description, OCR/text, number-color extraction, color grouping, chess, music, geometry, and math worksheets.

Table 18: Prompt contracts used by the live staged planner and the vision-tool backend.

K.1 System JSON Guard

For remote API backends, the runner prepends a provider-compatible system instruction whose invariant contract is:

```
[JSON OUTPUT REQUIRED]
You are a JSON-only assistant. Your ENTIRE response must be a valid JSON object.
CRITICAL RULES:
1. Response STARTS with { and ENDS with }.
2. NO text before or after the JSON object.
3. NO markdown formatting.
4. NO thinking tags such as <think> or <analysis>.
5. Valid JSON syntax only.
Output a single JSON object now:
```

Dataset-language suffixes may be appended for cross-lingual records, but they do not alter the JSON-only contract.

K.2 Stage 1: Abstract Planning

Stage 1 is tool-agnostic. The model sees the user request and any record-level extra instruction, but it does not see concrete runtime tool names. The prompt asks for a complexity-appropriate plan, concrete task-specific labels, and a matching abstract dependency graph:

```
[ABSTRACT PLANNING TASK]
```



```

776
777 ## User Request
778 {user_query}{extra_block}
779
780 ## Task
781 Create a high-level abstract plan to solve THIS SPECIFIC REQUEST.
782 Also express the same planning intent as an abstract dependency DAG.
783
784 CRITICAL RULES:
785 1. Length depends on complexity; do not force a fixed number of steps.
786 2. Use concrete terms from the request, not placeholders.
787 3. Do not name executable tool IDs or APIs.
788 4. Keep 'steps' and 'abs_plan_dag.nodes' aligned.
789 5. Add an edge only when the target step needs the source step's fact,
790    output, or established state.
791 6. Expose independent branches when steps can run independently.
792
793 Required JSON shape:
794 {
795   "steps": [{"step_index": 0, "description": "..."}],
796   "abs_plan_dag": {
797     "nodes": [{"node_id": "a0", "step_index": 0, "label": "...",
798               "step_type": "thought", "needs_tool": false}],
799     "edges": [{"source": "a0", "target": "a1", "edge_type": "dependency"}]
800   },
801   "valid_execution_order": ["a0", "a1"],
802   "reasoning": "..."
803 }

```

804 The normalized Stage-1 graph $\hat{G}_{\text{plan}}$ is stored separately from the concrete Stage-3 execution plan
805 and is the default source for planning-intent evaluation.

806 K.3 Stage 2: Tool Sufficiency

807 Stage 2 receives the abstract plan and the visible tool inventory. In the default record-level tool scope
808 this inventory comes from the task record; in global-scope ablations it comes from the executable
809 runtime tool library. The prompt is intentionally conservative:

```

810 [TOOL CREATION TASK]
811
812 ## User Request
813 {user_query}
814
815 ## Abstract Plan
816 {plan_block}
817
818 ## Existing Tools
819 {tools_block}
820
821 ## Task
822 Analyze if the Existing Tools are sufficient to execute the Abstract Plan.
823
824 IMPORTANT:
825 1. Check each abstract step against the existing tools.
826 2. If all steps can be handled, return an empty "new_tools" list.
827 3. Propose a new tool only for a clear capability gap.
828 4. Do not create duplicate tools with slightly different names.
829
830 Required JSON shape:
831 {
832   "reasoning": "<tool sufficiency analysis>",
833   "new_tools": []
834 }

```

835 Only Stage-2 suggestions that map to executable runtime tools are merged into the answer-mode tool
836 list; non-executable suggestions are logged but not exposed as callable tools.

837 K.4 Stage 3: Concrete Execution Planning

838 Stage 3 maps the Stage-1 abstract plan into a concrete execution graph for tool execution. In answer
839 mode, the prompt uses the same executable tool context that Stage 4 will later expose, so planned
840 tool choices are checked against the runtime tool universe.

```
841 [EXECUTION PLANNING TASK]
842
843 ## User Request
844 {user_query}
845
846 {attachment_block}{tool_creation_bridge_block}
847 ## Proposed Abstract Plan (Reference)
848 {plan_block}
849
850 ## Available Tools (Existing + Created)
851 {tools_block}
852
853 ## Task
854 Create a CONCRETE execution plan as a Directed Acyclic Graph (DAG).
855 Follow the Abstract Plan logic where possible, but map it to specific tools.
856
857 File and dependency rules:
858 - File-reading tools must provide the exact 'file_path' from attachments.
859 - Audio attachments should start with 'audio_transcription'.
860 - If step B uses step A's output, add a data dependency edge A -> B.
861 - Use <nX> in arguments to reference output from node nX.
862 - Do not add 'submit_final_answer' as a Stage-3 node or tool call.
863
864 Required JSON shape:
865 {
866   "plan_dag": {
867     "nodes": [{ "node_id": "n0", "step_index": 0, "label": "...",
868                 "step_type": "tool", "tool_id": "<TOOL_ID>",
869                 "arguments": { "ARG_NAME": "<ARG_VALUE>" },
870                 "output_vars": [ "<n0>" ], "needs_tool": true }],
871     "edges": [ { "source": "n0", "target": "n1", "edge_type": "data_dep" } ]
872   },
873   "tool_calls": [ { "call_index": 0, "node_id": "n0", "tool_id": "<TOOL_ID>",
874                   "arguments": [ { "name": "<ARG_NAME>", "value": "<ARG_VALUE>" } ] } ],
875   "final_answer": { "answer_type": "string", "answer": null }
876 }
```

877 The attachment block lists available local files and paths. For the async GAIA variant, an additional
878 prompt block asks the model to preserve independent branches and avoid serializing steps without
879 dependency justification.

880 K.5 Stage 4: Iterative Execution and Repair

881 Answer-mode execution is a ReAct-style loop. At every turn, the model sees the current task, the
882 Stage-3 plan when available, recent observations, and the visible runtime tools. The decision prompt
883 is regenerated after every observation:

```
884 Task:
885 {user_query}
886
887 {attachment_text}
888 {plan_text}
889
890 {previous_steps_and_observations_or_first_step}
891 {obs_text}
892 {reflection_text}
893 {image_policy_text}
894 {document_policy_text}
895 Available tools:
896 {available_tools_str}
897
```

898 Return EXACTLY ONE JSON object for the next action:

```
899 {
900   "tool_id": "tool_name",
901   "arguments": {"param": "value"},
902   "reasoning": "brief reason for this one action"
903 }
```

904

905 Rules:

- 906 - Choose exactly one next tool action.
- 907 - Do not output a multi-step plan, 'plan_dag', 'tool_calls', or 'final_answer'.
- 908 - If a recent observation supports one clear answer, call 'submit_final_answer'.
- 909 - Do not repeat the same tool call with the same arguments.
- 910 - For spreadsheets or tables, inspect columns/head/shape before computing.

911 Dynamic guidance is added only when relevant: image attachments trigger either a native-vision
912 policy or an image-recognition-first policy; audio files trigger transcription guidance; document-heavy
913 tasks trigger an attachment-first policy; and repeated low-yield actions trigger loop-control notes.
914 If the response is malformed, the repair prompt appends the failed output and asks for exactly one
915 corrected tool action with arguments and reasoning. A separate argument-refinement prompt is used
916 when the selected tool is valid but the parameter names do not match the tool signature.

917 K.6 Stage 5: Final Answer Synthesis

918 After iterative execution, Stage 5 summarizes the bounded observation history into a structured final
919 answer:

920 [ANSWER GENERATION]

921

922 ## User Request

923 {user_query}

924

925 ## Tool Execution History

926 {obs_text}

927

928 ## Task

929 Synthesize the Tool Execution History to provide the Final Answer to the User Request.

930 Strictly output the answer in JSON format.

931

932 Required JSON shape:

```
933 {
934   "final_answer": {
935     "answer": "The final concise answer string",
936     "reasoning": "Brief justification based on tool outputs"
937   }
938 }
```

939 The runner records a Stage-5 synthesis, but the answer-review sheet uses the Stage-4
940 submit_final_answer field as the judged answer. The Stage-5 answer is retained as pre-final
941 synthesis and audit context and does not override a Stage-4 submission.

942 K.7 Vision Tool Backend

943 The image-recognition runtime tool is itself backed by a structured vision prompt. The vision model
944 is instructed to act as an evidence extractor for the planner, not as the final task solver:

945 You are a vision analysis tool being called by a planning LLM.

946 Your role is to extract ACCURATE, COMPLETE, and STRUCTURED information from images.

947

948 CRITICAL RULES:

- 949 1. Extract all relevant information systematically.
- 950 2. Preserve exact values: numbers, text, colors, and positions.
- 951 3. Use clear, structured output that the planner can parse.
- 952 4. If uncertain, state what you see without guessing.
- 953 5. Focus on extracting data, not solving the task.

954 The selected task argument appends a task-specific extractor request. The implemented families
955 include general description, OCR/text extraction, number extraction, color-conditioned extraction,
956 chess, music-sheet, geometry, and math-worksheet variants. For number/color extraction, local OCR
957 and color candidates may be provided as a scaffold; the vision model is then asked to verify, correct,
958 and complete the candidate list against the image itself.

To test whether the augmentation effect and three-facet separation generalize beyond open-weight models, we evaluate eight API-dependent closed models from the OpenAI and Gemini families (*GPT-5.5*, *GPT-5.4*, *GPT-5.4-Mini*, *GPT-5.4-Nano*, *Gemini 3 Flash*, *Gemini 3.1 Flash-Lite*, *Gemini 2.5 Flash*, and *Gemini 2.5 Flash-Lite*) on the same 165 GAIA tasks under the same pipeline and metric definitions used in the main experiments. Table 19 reports the combined headline results, Table 20 reports the multi-order score lift, and Table 21 reports the tool-clean planning-answer diagnostic.

The closed-model runs use the same Augmented GAIA reference convention as the main text: native chain references plus Gemma 4 behavior-preserving replay-filtered async orderings. When a Stage-1 abstract planning graph is empty, the closed-model analysis applies the same Stage-3 concrete-plan fallback for planning metrics only; tool traces and submitted-answer fields are unchanged. The augmentation lift is positive for all eight closed models on both multi-order subsets (8/8), extending the main-text finding from 7/7 open-weight models to 15/15 models total. The three-facet separation is also pronounced: *GPT-5.5* attains the highest Accepted Answer rate in the full pool (0.673), while *Gemini 3 Flash* reaches the strongest Gemini-family Accepted Answer rate (0.552); both remain comparable to the open-weight pool on structural planning metrics rather than uniformly dominating it. On the planning-answer diagnostic, the closed-model pool mean Δr is -0.004 (versus $+0.014$ for the open-weight pool), with model-level CIs spanning zero in both directions. Closed models are reported separately because their weights are not publicly available and full reproduction depends on API access; the open-weight evaluation in the main text remains the primary reproducible result.

Model	Params	(A) Plan Structural Fidelity			(B) Tool Invocation Accuracy			(C) End-Task Answer Correctness			
		NODEF1 (Native = Aug.)	EDGEF1		TOOLF1	PARAMF1	PARAMVALUEF1	EM	TOKENF1	Raw rate	Accepted*
			Native	Aug. Δ							
Closed-model											
GPT-5.5	–	0.523	0.097	0.110 +0.012	0.472	0.445	0.262	0.582	0.617	0.582	0.673
GPT-5.4	–	0.535	0.098	0.110 +0.012	0.549	0.538	0.264	0.267	0.302	0.212	0.352
GPT-5.4-Mini	–	0.509	0.119	0.130 +0.011	0.623	0.564	0.246	0.115	0.170	0.109	0.158
GPT-5.4-Nano	–	0.475	0.062	0.071 +0.009	0.558	0.543	0.220	0.127	0.171	0.103	0.170
Gemini 3 Flash	–	0.467	0.091	0.100 +0.009	0.525	0.557	0.272	0.497	0.545	0.467	0.552
Gemini 3.1 Flash-Lite	–	0.442	0.093	0.098 +0.005	0.550	0.560	0.257	0.303	0.363	0.321	0.406
Gemini 2.5 Flash	–	0.526	0.118	0.136 +0.018	0.438	0.466	0.241	0.200	0.239	0.194	0.321
Gemini 2.5 Flash-Lite	–	0.497	0.100	0.110 +0.010	0.528	0.531	0.227	0.182	0.213	0.182	0.248
Closed-model mean	–	0.497	0.097	0.108 +0.011	0.530	0.526	0.249	0.284	0.328	0.271	0.360
Open-weight											
Mistral Large 3	675B	0.514	0.101	0.107 +0.006	0.412	0.458	0.206	0.182	0.223	0.224	0.194
Llama-3.1	405B	0.519	0.107	0.118 +0.011	0.473	0.518	0.279	0.176	0.217	0.261	0.182
Llama-3.3	70B	0.486	0.091	0.105 +0.014	0.527	0.572	0.232	0.182	0.232	0.219	0.206
Gemma 4	31B	0.519	0.101	0.116 +0.015	0.537	0.579	0.264	0.339	0.384	0.445	0.424
Mistral Small 3.2	24B	0.512	0.096	0.109 +0.013	0.545	0.569	0.233	0.103	0.148	0.179	0.145
Llama-4 Maverick	17B	0.494	0.092	0.105 +0.013	0.424	0.454	0.201	0.218	0.263	0.294	0.261
Gemma 3	12B	0.499	0.079	0.091 +0.012	0.549	0.567	0.241	0.091	0.136	0.158	0.115
Open-weight mean	–	0.506	0.095	0.107 +0.012	0.495	0.531	0.237	0.184	0.229	0.254	0.218

Table 19: Combined three-facet headline results for the eight API-dependent closed models and the seven open-weight models. Open-weight rows are copied from Table 3; closed-model rows follow the same submitted-answer convention and Augmented GAIA reference set.

Subset (n)	Metric	Native	Augmented	Δ	Δ (%)	95% CI of Δ	Models with $\Delta > 0$
Closed-model							
GAIA Level 2* ($n = 29$)	NODEF1	0.499	0.499	+0.000	+0.0%	— (by construction)	0/8
	EDGEF1	0.086	0.114	+0.028	+32.6%	[+0.019, +0.037]	8/8
	PLANNINGSCORE	0.286	0.300	+0.014	+4.9%	[+0.010, +0.019]	8/8
All* ($n = 54$)	NODEF1	0.491	0.491	+0.000	+0.0%	— (by construction)	0/8
	EDGEF1	0.074	0.102	+0.029	+38.6%	[+0.023, +0.035]	8/8
	PLANNINGSCORE	0.278	0.293	+0.014	+5.1%	[+0.011, +0.018]	8/8
Open-weight							
GAIA Level 2* ($n = 29$)	NODEF1	0.539	0.539	+0.000	+0.0%	— (by construction)	0/7
	EDGEF1	0.088	0.124	+0.036	+40.6%	[+0.0271, +0.0439]	7/7
	PLANNINGSCORE	0.313	0.331	+0.018	+5.7%	[+0.0135, +0.0219]	7/7
All* ($n = 54$)	NODEF1	0.515	0.515	+0.000	+0.0%	— (by construction)	0/7
	EDGEF1	0.074	0.107	+0.033	+45.4%	[+0.0263, +0.0398]	7/7
	PLANNINGSCORE	0.294	0.311	+0.017	+5.7%	[+0.0132, +0.0199]	7/7

* Multi-order-rich denotes tasks with at least two retained async orderings beyond the native chain.

Table 20: Multi-order score lift for closed-model and open-weight pools. Bootstrap intervals use the same cross-model resampling protocol as Table 4.

Model	Params	N	Accepted rate	r_{native}	r_{aug}	Δr	95% CI of Δr
Closed-model							
<i>GPT-5.5</i>	—	56	0.643	+0.311	+0.322	+0.010	[−0.022, +0.040]
<i>GPT-5.4</i>	—	62	0.355	−0.008	−0.032	−0.025	[−0.075, +0.027]
<i>GPT-5.4-Mini</i>	—	52	0.154	−0.059	−0.076	−0.017	[−0.042, +0.000]
<i>GPT-5.4-Nano</i>	—	61	0.148	+0.183	+0.196	+0.013	[−0.019, +0.057]
<i>Gemini 3 Flash</i>	—	53	0.434	+0.317	+0.308	−0.009	[−0.030, +0.000]
<i>Gemini 3.1 Flash-Lite</i>	—	52	0.288	−0.013	−0.023	−0.010	[−0.033, +0.000]
<i>Gemini 2.5 Flash</i>	—	52	0.250	−0.075	−0.050	+0.026	[−0.020, +0.077]
<i>Gemini 2.5 Flash-Lite</i>	—	46	0.130	+0.361	+0.342	−0.019	[−0.044, +0.000]
Closed-model mean / total	—	434	0.300	+0.127	+0.123	−0.004	—
Open-weight							
<i>Mistral Large 3</i>	675B	53	0.170	−0.156	−0.141	+0.014	[−0.025, +0.074]
<i>Llama-3.1</i>	405B	49	0.184	+0.217	+0.207	−0.010	[−0.034, +0.000]
<i>Llama-3.3</i>	70B	58	0.172	−0.080	−0.037	+0.043	[−0.022, +0.145]
<i>Gemma 4</i>	31B	62	0.371	+0.258	+0.276	+0.018	[−0.007, +0.053]
<i>Mistral Small 3.2</i>	24B	60	0.150	+0.039	+0.056	+0.016	[−0.039, +0.092]
<i>Llama-4 Maverick</i>	17B	59	0.254	+0.045	+0.079	+0.034	[−0.027, +0.110]
<i>Gemma 3</i>	12B	57	0.088	+0.042	+0.027	−0.015	[−0.046, −0.000]
Open-weight mean / total	—	398	0.201	+0.052	+0.066	+0.014	—

Table 21: Tool-clean planning–answer diagnostic for closed-model and open-weight pools on the GAIA Level 2 *Text/Document/Audio* slice.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes].

Justification: The abstract and introduction state the three-facet evaluation contribution, the GAIA-derived benchmark construction, and the empirical scope; the corresponding evidence appears in Sections 4.2–5 and Appendix L.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: Section 6 discusses single-model replay filtering, statistical limits from the 165-task and 7-model open-weight pool, domain imbalance, and cross-benchmark scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A].

Justification: The paper proposes and evaluates a benchmark and metric protocol; it does not contain theoretical theorems or formal proof claims.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes].

Justification: Sections 5 and Appendices A, J, and K specify the pipeline, data construction, model pool, bootstrap protocol, prompts, and evaluation settings. The anonymized artifacts provide derived annotations, rebuild scripts, and Croissant metadata; GAIA raw questions, final answers, and attachments must be obtained from the official gated source rather than redistributed in the release. Appendix L separately reports the API-dependent closed-model extension.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility.

1087 In the case of closed-source models, it may be that access to the model is limited in
1088 some way (e.g., to registered users), but it should be possible for other researchers
1089 to have some path to reproducing or verifying the results.

1090 5. Open access to data and code

1091 Question: Does the paper provide open access to the data and code, with sufficient instruc-
1092 tions to faithfully reproduce the main experimental results, as described in supplemental
1093 material?

1094 Answer: [Yes].

1095 Justification: The anonymized release contains the evaluation code, data-preparation scripts,
1096 Croissant metadata, sanitized result summaries, and derived Augmented GAIA annotations
1097 needed to reproduce the open-weight experiments. To respect GAIA’s gated access terms,
1098 the release does not contain GAIA raw questions, final answers, or attachments; users rebuild
1099 the local data layout by combining the official GAIA source with the released annotations.
1100 Closed-model results are reported separately because reproduction depends on API access.

1101 Guidelines:

- 1102 • The answer [N/A] means that paper does not include experiments requiring code.
- 1103 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/](https://neurips.cc/public/guides/CodeSubmissionPolicy)
1104 [public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 1105 • While we encourage the release of code and data, we understand that this might not
1106 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1107 including code, unless this is central to the contribution (e.g., for a new open-source
1108 benchmark).
- 1109 • The instructions should contain the exact command and environment needed to run to
1110 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1111 • The authors should provide instructions on data access and preparation, including how
1112 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1113 • The authors should provide scripts to reproduce all experimental results for the new
1114 proposed method and baselines. If only a subset of experiments are reproducible, they
1115 should state which ones are omitted from the script and why.
- 1116 • At submission time, to preserve anonymity, the authors should release anonymized
1117 versions (if applicable).
- 1118 • Providing as much information as possible in supplemental material (appended to the
1119 paper) is recommended, but including URLs to data and code is permitted.
- 1120

1121 6. Experimental setting/details

1122 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1123 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1124 Answer: [Yes].

1125 Justification: Section 5 gives the experimental setting, including seven open-weight models
1126 in the main evaluation, while Appendix L reports the eight API-dependent closed models.
1127 Appendices A, H, and K give environment, filtering, bootstrap, and prompt details.

1128 Guidelines:

- 1129 • The answer [N/A] means that the paper does not include experiments.
- 1130 • The experimental setting should be presented in the core of the paper to a level of detail
1131 that is necessary to appreciate the results and make sense of them.
- 1132 • The full details can be provided either with the code, in appendix, or as supplemental
1133 material.

1134 7. Experiment statistical significance

1135 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1136 information about the statistical significance of the experiments?

1137 Answer: [Yes].

Justification: The main score-lift and planning–answer tables report bootstrap confidence intervals, and Appendix A.3 defines the resampling protocols.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes].

Justification: Appendix A.4 reports the software environment, CPU/GPU hardware, RAM, inference-only setup, and the distinction between endpoint generation time and cached metric recomputation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes].

Justification: The work is an evaluation benchmark built from existing benchmarks and derived annotations, preserves anonymity, does not collect human-subject data, documents external asset terms in Appendix A.5, and avoids redistributing GAIA raw validation/test data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.

- 1191 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
1192 eration due to laws or regulations in their jurisdiction).

1193 10. Broader impacts

1194 Question: Does the paper discuss both potential positive societal impacts and negative
1195 societal impacts of the work performed?

1196 Answer: [No].

1197 Justification: The paper focuses on diagnostic benefits and limitations of separated agent
1198 evaluation but does not include a dedicated broader-impact section discussing positive and
1199 negative societal impacts in detail.

1200 Guidelines:

- 1201 • The answer [N/A] means that there is no societal impact of the work performed.
- 1202 • If the authors answer [N/A] or [No], they should explain why their work has no societal
1203 impact or why the paper does not address societal impact.
- 1204 • Examples of negative societal impacts include potential malicious or unintended uses
1205 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
1206 (e.g., deployment of technologies that could make decisions that unfairly impact specific
1207 groups), privacy considerations, and security considerations.
- 1208 • The conference expects that many papers will be foundational research and not tied
1209 to particular applications, let alone deployments. However, if there is a direct path to
1210 any negative applications, the authors should point it out. For example, it is legitimate
1211 to point out that an improvement in the quality of generative models could be used to
1212 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
1213 that a generic algorithm for optimizing neural networks could enable people to train
1214 models that generate Deepfakes faster.
- 1215 • The authors should consider possible harms that could arise when the technology is
1216 being used as intended and functioning correctly, harms that could arise when the
1217 technology is being used as intended but gives incorrect results, and harms following
1218 from (intentional or unintentional) misuse of the technology.
- 1219 • If there are negative societal impacts, the authors could also discuss possible mitigation
1220 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1221 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1222 feedback over time, improving the efficiency and accessibility of ML).

1223 11. Safeguards

1224 Question: Does the paper describe safeguards that have been put in place for responsible
1225 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1226 image generators, or scraped datasets)?

1227 Answer: [N/A].

1228 Justification: The paper does not release pretrained model weights, image generators, or
1229 a scraped dataset with high misuse risk. The released artifact contains derived evaluation
1230 annotations, sanitized aggregate summaries, code, checksums, and Croissant metadata; it
1231 does not redistribute GAIA raw questions, final answers, or attachments.

1232 Guidelines:

- 1233 • The answer [N/A] means that the paper poses no such risks.
- 1234 • Released models that have a high risk for misuse or dual-use should be released with
1235 necessary safeguards to allow for controlled use of the model, for example by requiring
1236 that users adhere to usage guidelines or restrictions to access the model or implementing
1237 safety filters.
- 1238 • Datasets that have been scraped from the Internet could pose safety risks. The authors
1239 should describe how they avoided releasing unsafe images.
- 1240 • We recognize that providing effective safeguards is challenging, and many papers do
1241 not require this, but we encourage authors to take this into account and make a best
1242 faith effort.

1243 12. Licenses for existing assets

1244 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1245 the paper, properly credited and are the license and terms of use explicitly mentioned and
1246 properly respected?

1247 Answer: [Yes].

1248 Justification: Source benchmarks and model families are cited in the paper, and Ap-
1249 pendix A.5 lists the external assets, source URLs where applicable, and the requirement to
1250 follow the original licenses or provider terms. The release respects GAIA's gated dataset
1251 terms by providing derived annotations and rebuild scripts rather than redistributing GAIA
1252 raw questions, final answers, or attachments.

1253 Guidelines:

- 1254 • The answer [N/A] means that the paper does not use existing assets.
- 1255 • The authors should cite the original paper that produced the code package or dataset.
- 1256 • The authors should state which version of the asset is used and, if possible, include a
1257 URL.
- 1258 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1259 • For scraped data from a particular source (e.g., website), the copyright and terms of
1260 service of that source should be provided.
- 1261 • If assets are released, the license, copyright information, and terms of use in the
1262 package should be provided. For popular datasets, `paperswithcode.com/datasets`
1263 has curated licenses for some datasets. Their licensing guide can help determine the
1264 license of a dataset.
- 1265 • For existing datasets that are re-packaged, both the original license and the license of
1266 the derived asset (if it has changed) should be provided.
- 1267 • If this information is not available online, the authors are encouraged to reach out to
1268 the asset's creators.

1269 13. New assets

1270 Question: Are new assets introduced in the paper well documented and is the documentation
1271 provided alongside the assets?

1272 Answer: [Yes].

1273 Justification: The paper documents the new benchmark annotations, tool layer, order augmen-
1274 tation, validation procedure, and prompt contracts in Sections 4.2–4.4 and Appendices C, J,
1275 and K. The accompanying anonymized dataset artifact documents the controlled annotation
1276 bundle, rebuild scripts, checksums, sanitized result summaries, and Croissant metadata.

1277 Guidelines:

- 1278 • The answer [N/A] means that the paper does not release new assets.
- 1279 • Researchers should communicate the details of the dataset/code/model as part of their
1280 submissions via structured templates. This includes details about training, license,
1281 limitations, etc.
- 1282 • The paper should discuss whether and how consent was obtained from people whose
1283 asset is used.
- 1284 • At submission time, remember to anonymize your assets (if applicable). You can either
1285 create an anonymized URL or include an anonymized zip file.

1286 14. Crowdsourcing and research with human subjects

1287 Question: For crowdsourcing experiments and research with human subjects, does the paper
1288 include the full text of instructions given to participants and screenshots, if applicable, as
1289 well as details about compensation (if any)?

1290 Answer: [N/A].

1291 Justification: The work does not involve crowdsourcing or a human-subject study; hu-
1292 man checks are internal audits of dependency annotations and model outputs rather than
1293 experiments on participants.

1294 Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A].

Justification: The paper does not conduct a human-subject study and does not collect participant data, so IRB approval is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes].

Justification: The paper documents LLM use in the core methodology: *GPT-4o* is used for dependency annotation and Gemma 4 is used as the replay model for behavior-preserving filtering; the 7 open-weight and 8 API-dependent closed evaluated models are listed as experimental subjects in Sections 5 and Appendix L.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.